

Tensorflow setup

```
In [ ]:  
# Python ≥3.5 is required  
import sys  
assert sys.version_info >= (3, 5)  
  
# Scikit-Learn ≥0.20 is required  
import sklearn  
assert sklearn.__version__ >= "0.20"  
  
try:  
    # %tensorflow_version only exists in Colab.  
    %tensorflow_version 2.x  
except Exception:  
    pass  
  
# TensorFlow ≥2.0 is required  
import tensorflow as tf  
assert tf.__version__ >= "2.0"  
  
# Common imports  
import numpy as np  
import os  
  
# to make this notebook's output stable across runs  
np.random.seed(42)  
  
# To plot pretty figures  
%matplotlib inline  
import matplotlib as mpl  
import matplotlib.pyplot as plt  
mpl.rcParams['axes'].labelsize=14  
mpl.rcParams['xtick'].labelsize=12  
mpl.rcParams['ytick'].labelsize=12  
  
# Ignore useless warnings (see SciPy issue #5998)  
import warnings  
warnings.filterwarnings(action="ignore", message="^internal gelsd")
```

Every import needed

```
In [ ]:  
from tensorflow import keras  
from keras import layers  
from tensorflow.keras.preprocessing import image_dataset_from_directory  
from tensorflow import keras  
from tensorflow.keras.preprocessing import image_dataset_from_directory  
  
from sklearn.metrics import classification_report  
from sklearn.metrics import confusion_matrix  
from sklearn.metrics import precision_recall_fscore_support  
from sklearn.metrics import accuracy_score  
from sklearn.metrics import f1_score  
  
import seaborn as sns  
  
import os  
import pathlib  
import time  
  
# Convert dataset to numpy function  
from tensorflow.data import Dataset  
  
from sklearn.metrics import precision_score
```

```
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import balanced_accuracy_score
from tensorflow.keras.utils import to_categorical
from itertools import cycle
from sklearn.metrics import auc

import pandas as pd
```

Load images

```
In [ ]: PATH = '.\\..\\dataset\\micro_expressions'

train_dir = os.path.join(PATH, 'train')
test_dir = os.path.join(PATH, 'test')

print(train_dir)
print(test_dir)
```

```
.\\..\\dataset\\micro_expressions\\train
.\\..\\dataset\\micro_expressions\\test
```

```
In [ ]: import pathlib

train_dir = pathlib.Path(train_dir)
test_dir = pathlib.Path(test_dir)

image_count = len(list(train_dir.glob('/*/*.jpg')))
print(image_count)

image_count = len(list(test_dir.glob('/*/*.jpg')))
print(image_count)
```

```
6938
1700
```

```
In [ ]: batch_size = 32
IMG_SIZE = (80, 80)

train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    validation_split=0.15,
    subset="training",
    seed=123,
    shuffle=True,
    image_size=IMG_SIZE,
    batch_size=batch_size)

val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    validation_split=0.15,
    subset="validation",
    seed=123,
    shuffle=True,
    image_size=IMG_SIZE,
    batch_size=batch_size)

test_ds = tf.keras.preprocessing.image_dataset_from_directory(
    test_dir,
    validation_split=None,
    seed=123,
```

```

shuffle=True,
image_size=IMG_SIZE,
batch_size=batch_size)

class_names = train_ds.class_names

train_ds = train_ds.cache().prefetch(1)
val_ds = val_ds.cache().prefetch(1)
test_ds = test_ds.cache().prefetch(1)

normalization_layer = keras.layers.experimental.preprocessing.Rescaling(1./255)

train_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
val_ds = val_ds.map(lambda x, y: (normalization_layer(x), y))
test_ds = test_ds.map(lambda x, y: (normalization_layer(x), y))

```

Found 6938 files belonging to 7 classes.
Using 5898 files for training.
Found 6938 files belonging to 7 classes.
Using 1040 files for validation.
Found 1700 files belonging to 7 classes.

Util functions used in the code

In []:

```

# Convert dataset to numpy
def batch_dataset_to_numpy(ds):
    in_data = []
    out_data = []
    for inp, out in ds:

        array = inp.numpy()
        for x in array:
            in_data.append(x)

        array = out.numpy()
        for x in array:
            out_data.append(x)

    return in_data, out_data

X_train, y_train = batch_dataset_to_numpy(train_ds)
X_val, y_val = batch_dataset_to_numpy(val_ds)
X_test, y_test = batch_dataset_to_numpy(test_ds)

```

In []:

```

def print_metrics_by_class(array, title=None, is_rate=False):
    i = 0
    print(title)

    for k in class_names:
        if is_rate == True:
            print('\t%s -> %.3f%%' %(k, array[i]))
        else:
            print('\t%s -> %.0f' %(k, array[i]))

        i = i + 1

def print_confusion_matrix_metrics(cnf_matrix):
    FP = cnf_matrix.sum(axis=0) - np.diag(cnf_matrix)
    FN = cnf_matrix.sum(axis=1) - np.diag(cnf_matrix)
    TP = np.diag(cnf_matrix)
    TN = cnf_matrix.sum() - (FP + FN + TP)

    FP = FP.astype(float)
    FN = FN.astype(float)
    TP = TP.astype(float)

```

```

TN = TN.astype(float)

print_metrics_by_class(FP, '\nFalse Positives:\n')
print_metrics_by_class(FN, '\nFalse Negatives:\n')
print_metrics_by_class(TP, '\nTrue Positives:\n')
print_metrics_by_class(TN, '\nTrue Negatives:\n')

# Sensitivity, hit rate, recall, or true positive rate
TPR = TP/(TP+FN)
print_metrics_by_class(TPR, '\nTrue Positive Rate:\n', is_rate=True)

# Specificity or true negative rate
TNR = TN/(TN+FP)
print_metrics_by_class(TNR, '\nTrue Negative Rate:\n', is_rate=True)

# Precision or positive predictive value
PPV = TP/(TP+FP)
print_metrics_by_class(PPV, '\nPositive Predictive Value:\n')

# Negative predictive value
NPV = TN/(TN+FN)
print_metrics_by_class(NPV, '\nNegative Predictive Value:\n')

# Fall out or false positive rate
FPR = FP/(FP+TN)
print_metrics_by_class(FPR, '\nFalse Positive Rate:\n', is_rate=True)

# False negative rate
FNR = FN/(TP+FN)
print_metrics_by_class(FNR, '\nFalse Negative Rate:\n', is_rate=True)

# False discovery rate
FDR = FP/(TP+FP)
print_metrics_by_class(FDR, '\nFalse Discovery Rate:\n', is_rate=True)

# Overall accuracy
ACC = (TP+TN)/(TP+FP+FN+TN)
print_metrics_by_class(ACC, '\nAccuracy:\n', is_rate=True)

def plot_confusion_matrix(y_pred, y_test, labels, normalize_m=None, return_m=False):

    fig, ax = plt.subplots(figsize=(10,10))
    matrix = confusion_matrix(y_test, tf.argmax(y_pred, axis=-1), normalize=normalize_m)

    fmt_m = '2'
    if normalize_m != None:
        fmt_m = '.3f'

    sns.heatmap(
        matrix,
        annot=True,
        fmt=fmt_m,
        cmap='Blues',
        cbar=False,
        xticklabels=class_names,
        yticklabels=class_names,
        ax=ax
    )

    plt.title('Confusion Matrix', fontsize = 20)
    plt.ylabel('Real Label', fontsize = 15)
    plt.xlabel('Predicted Label', fontsize = 15)
    plt.show()

    if return_m == True:
        return matrix

def plot_loss_and_accuracy(history):

```

```

history_data_frame = pd.DataFrame(history.history)

print(history_data_frame.loc[:, [ 'accuracy']].max(), history_data_frame.loc[:, [ 'val_
```

`plt.figure(figsize=(15,10))
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, linestyle='-', marker='.', color='b', label="Training loss")
plt.plot(epochs, val_loss, linestyle='-', marker='.', color='r', label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show();`

```

plt.clf()
plt.figure(figsize=(15,10))
acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]
plt.plot(epochs, acc, linestyle='-', marker='.', color='b', label="Training accuracy")
plt.plot(epochs, val_acc, linestyle='-', marker='.', color='r', label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.grid(True)
plt.gca().set_ylim(0, 1) # set the vertical range to [0-1]
plt.show()
```

```

def print_dataset_class_balance(class_names, y, title=None):
    print(title)
    y_np = np.array(y)
    length = len(y_np)

    for i, x in enumerate(class_names):
        count = len(y_np[y_np == i])
        print(f" {class_names[i]} = {(count/length):.2f} %")

    print('\n')
```

```

def batch_dataset_to_numpy(ds):
    in_data = []
    out_data = []
    for inp, out in ds:

        array = inp.numpy()
        for x in array:
            in_data.append(x)

        array = out.numpy()
        for x in array:
            out_data.append(x)

    return in_data, out_data
```

```

def print_dataset_information():
    # Cardinality (number of batches)
    print('Train Cardinality (number of batches): ', train_ds.cardinality().numpy())
    print('Validation cardinality (number of batches): ', val_ds.cardinality().numpy())
    print('Test cardinality (number of batches): ', test_ds.cardinality().numpy(), end='\n')

    print('Class names: ', class_names, end='\n\n')

    print('Shape of X_train: ', np.array(X_train).shape)
```

```

print('Shape of y_train: ', np.array(y_train).shape, end='\n\n')
print('Shape of X_val: ', np.array(X_val).shape)
print('Shape of y_val: ', np.array(y_val).shape, end='\n\n')
print('Shape of X_test: ', np.array(X_test).shape)
print('Shape of y_test: ', np.array(y_test).shape, end='\n\n')

print_dataset_class_balance(class_names, y_train, 'Class Balance - Train Data:\n')
print_dataset_class_balance(class_names, y_val, 'Class Balance - Validation Data:\n')
print_dataset_class_balance(class_names, y_test, 'Class Balance - Test Data:\n')

```

In []:

```

#https://scikit-learn.org/stable/auto_examples/model_selection/plot_roc.html#sphx-glr-auto-
graph_colors = ['aqua', 'darkorange', 'cornflowerblue', 'green', 'red', 'blue', 'grey']

def plot_multiclass_roc_curve(y_test, y_pred, class_names, single_class = None, title = None):
    if title == None:
        title = ''

    # Binarize the output
    y_b = to_categorical(y_test, num_classes = len(class_names))
    yp = y_pred

    # Compute ROC curve and ROC area for each class
    fpr = dict()
    tpr = dict()
    thresh = dict()
    roc_auc = dict()

    for i in range(len(class_names)):
        fpr[str(i)], tpr[str(i)], thresh[str(i)] = roc_curve(y_b[:, i], yp[:, i])
        roc_auc[str(i)] = auc(fpr[str(i)], tpr[str(i)])

    lw = 2 # Linewidth

    # Compute micro-average ROC curve and ROC area
    fpr['micro'], tpr['micro'], thresh['micro'] = roc_curve(y_b.ravel(), yp.ravel())
    roc_auc['micro'] = auc(fpr['micro'], tpr['micro'])

    # First aggregate all false positive rates
    all_fpr = np.unique(np.concatenate([fpr[str(i)] for i in range(len(class_names))]))

    # Then interpolate all ROC curves at this points
    mean_tpr = np.zeros_like(all_fpr)
    for i in range(len(class_names)):
        mean_tpr += np.interp(all_fpr, fpr[str(i)], tpr[str(i)])

    # Finally average it and compute AUC
    mean_tpr /= len(class_names)

    fpr['macro'] = all_fpr
    tpr['macro'] = mean_tpr
    roc_auc['macro'] = auc(fpr['macro'], tpr['macro'])

    # Plot all ROC curves
    plt.figure(figsize=(15,15))
    plt.plot(
        fpr['micro'],
        tpr['micro'],
        label = 'Micro-Avg (area = {0:0.2f})'.format(roc_auc['micro']),
        color = 'deeppink',
        linestyle = '--',
        linewidth = lw
    );
    plt.plot(

```

```

        fpr['macro'],
        tpr['macro'],
        label = 'Macro-Avg (area = {0:0.2f})'.format(roc_auc['macro']),
        color = 'navy',
        linestyle = '--',
        linewidth = lw
    );

    colors = cycle(graph_colors)
    for i, color in zip(range(len(class_names)), colors):
        if single_class == None:
            plt.plot(
                fpr[str(i)],
                tpr[str(i)],
                color = color,
                lw = lw,
                label = 'ROC - "{0}" (area = {1:0.2f})'.format(class_names[i].capitalize())
            );
        else:
            if single_class == i:
                plt.plot(
                    fpr[str(i)],
                    tpr[str(i)],
                    color = color,
                    lw = lw,
                    label = 'ROC - "{0}" (area = {1:0.2f})'.format(class_names[i].capitalize())
                );
            else:
                pass

plt.plot([0, 1],[0, 1], 'k--', lw = lw, label = 'Random Guessing')
plt.xlim([-0.025, 1.0])
plt.ylim([0.0, 1.025])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve (TPR vs FPR)' + title)
plt.legend(loc = 'lower right')
plt.show()

macro_roc_auc_ovo = roc_auc_score(y_b, yp, multi_class="ovo",
                                   average="macro")
weighted_roc_auc_ovo = roc_auc_score(y_b, yp, multi_class="ovo",
                                      average="weighted")

macro_roc_auc_ovr = roc_auc_score(y_b, yp, multi_class="ovr",
                                   average="macro")
weighted_roc_auc_ovr = roc_auc_score(y_b, yp, multi_class="ovr",
                                      average="weighted")

print("One-vs-One ROC AUC scores:\n\t{:.6f} (macro),\n\t{:.6f} "
      "(weighted by prevalence)"
      .format(macro_roc_auc_ovo, weighted_roc_auc_ovo))
print("One-vs-Rest ROC AUC scores:\n\t{:.6f} (macro),\n\t{:.6f} "
      "(weighted by prevalence)"
      .format(macro_roc_auc_ovr, weighted_roc_auc_ovr))

```

```

In [ ]: def print_general_metrics(y_pred, y_test):
    pred = tf.argmax(y_pred, axis = -1)
    precision, recall, fbeta_score, support = precision_recall_fscore_support(y_test, pred)

    print('\nClassification Report:\n')
    print(classification_report(y_test, pred), end='\n\n')

    print('Accuracy: %.3f\n' % accuracy_score(y_test, pred))
    print('Balanced accuracy: %.3f\n' % balanced_accuracy_score(y_test, pred))
    print('F1_Score: %.3f\n' % f1_score(y_test, pred, average='weighted'))

```

```
print('Precision: %.3f\n' % precision)
print('Recall: %.3f\n' % recall)
print('ROC AUC (OVR): %.3f\n' % roc_auc_score(y_test, y_pred, multi_class = 'ovr'))
print('ROC AUC (OVO): %.3f\n' % roc_auc_score(y_test, y_pred, multi_class = 'ovo'))
```

Dataset overview

```
In [ ]: print_dataset_information()

Train Cardinality (number of batches): 185
Validation cardinality (number of batches): 33
Test cardinality (number of batches): 54

Class names: ['anger', 'disgust', 'fear', 'happiness', 'neutral', 'sadness', 'surprise']

Shape of X_train: (5898, 80, 80, 3)
Shape of y_train: (5898,)

Shape of X_val: (1040, 80, 80, 3)
Shape of y_val: (1040,)

Shape of X_test: (1700, 80, 80, 3)
Shape of y_test: (1700,)

Class Balance - Train Data:

anger = 0.19 %
disgust = 0.08 %
fear = 0.06 %
happiness = 0.25 %
neutral = 0.08 %
sadness = 0.19 %
surprise = 0.14 %

Class Balance - Validation Data:

anger = 0.17 %
disgust = 0.09 %
fear = 0.06 %
happiness = 0.27 %
neutral = 0.08 %
sadness = 0.17 %
surprise = 0.14 %

Class Balance - Test Data:

anger = 0.19 %
disgust = 0.08 %
fear = 0.06 %
happiness = 0.25 %
neutral = 0.08 %
sadness = 0.19 %
surprise = 0.14 %
```

```
In [ ]: plt.figure(figsize=(15, 10))
for images, labels in train_ds.take(1):
    for i in range(10):
        ax = plt.subplot(2, 10, i + 1)
        plt.imshow(images[i].numpy().astype("float32"))
        plt.title(class_names[labels[i]])
        plt.axis("off")
```



Begginning of Tensorflow analysis

Base model without data augmentation

In []:

```

keras.backend.clear_session()
tf.random.set_seed(42)
np.random.seed(42)

inputs = keras.Input(shape=(80, 80, 3))

x = keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu', padding='same')(inputs)
x = keras.layers.MaxPooling2D()(x)

x = keras.layers.Conv2D(64, 3, activation='relu', padding='VALID')(x)
x = keras.layers.MaxPooling2D()(x)

x = keras.layers.Conv2D(120, 3, activation='relu', padding='VALID')(x)
x = keras.layers.MaxPooling2D()(x)

x = keras.layers.Flatten()(x)

x = keras.layers.Dense(90, activation='relu')(x)
x = keras.layers.Dropout(0.5)(x)
x = keras.layers.Dense(60, activation='relu')(x)
x = keras.layers.Dropout(0.5)(x)

outputs = keras.layers.Dense(7, activation="softmax")(x)

model = keras.Model(inputs=inputs, outputs=outputs)
model.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	[None, 80, 80, 3]	0
conv2d (Conv2D)	(None, 80, 80, 32)	896
max_pooling2d (MaxPooling2D)	(None, 40, 40, 32)	0
conv2d_1 (Conv2D)	(None, 38, 38, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 19, 19, 64)	0
conv2d_2 (Conv2D)	(None, 17, 17, 120)	69240
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 120)	0
flatten (Flatten)	(None, 7680)	0
dense (Dense)	(None, 90)	691290
dropout (Dropout)	(None, 90)	0
dense_1 (Dense)	(None, 60)	5460
dropout_1 (Dropout)	(None, 60)	0
dense_2 (Dense)	(None, 7)	427
<hr/>		
Total params: 785,809		

```
Trainable params: 785,809  
Non-trainable params: 0
```

```
In [ ]:  
L = keras.losses.SparseCategoricalCrossentropy()  
  
model.compile(loss=L, optimizer="adam", metrics=["accuracy"])
```

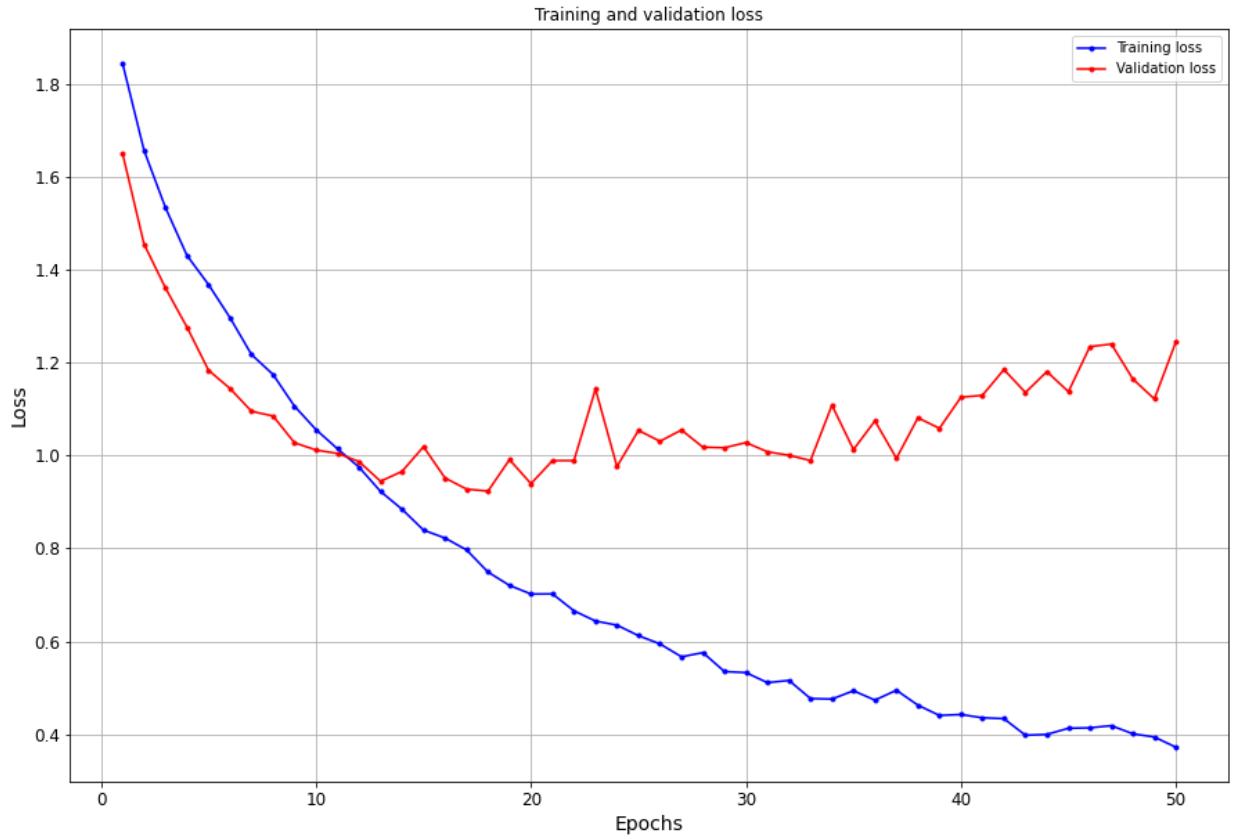
```
In [ ]:  
history = model.fit(  
    train_ds,  
    validation_data=val_ds,  
    epochs=50  
)  
  
Epoch 1/50  
185/185 [=====] - 2s 8ms/step - loss: 1.8441 - accuracy: 0.2567 -  
val_loss: 1.6502 - val_accuracy: 0.3683  
Epoch 2/50  
185/185 [=====] - 1s 7ms/step - loss: 1.6568 - accuracy: 0.3694 -  
val_loss: 1.4549 - val_accuracy: 0.4702  
Epoch 3/50  
185/185 [=====] - 1s 7ms/step - loss: 1.5331 - accuracy: 0.4262 -  
val_loss: 1.3602 - val_accuracy: 0.5192  
Epoch 4/50  
185/185 [=====] - 1s 7ms/step - loss: 1.4303 - accuracy: 0.4742 -  
val_loss: 1.2760 - val_accuracy: 0.5356  
Epoch 5/50  
185/185 [=====] - 1s 7ms/step - loss: 1.3679 - accuracy: 0.4941 -  
val_loss: 1.1834 - val_accuracy: 0.5683  
Epoch 6/50  
185/185 [=====] - 1s 7ms/step - loss: 1.2965 - accuracy: 0.5227 -  
val_loss: 1.1442 - val_accuracy: 0.5856  
Epoch 7/50  
185/185 [=====] - 1s 7ms/step - loss: 1.2172 - accuracy: 0.5605 -  
val_loss: 1.0952 - val_accuracy: 0.6029  
Epoch 8/50  
185/185 [=====] - 2s 9ms/step - loss: 1.1744 - accuracy: 0.5734 -  
val_loss: 1.0851 - val_accuracy: 0.6279  
Epoch 9/50  
185/185 [=====] - 2s 9ms/step - loss: 1.1061 - accuracy: 0.6063 -  
val_loss: 1.0274 - val_accuracy: 0.6433  
Epoch 10/50  
185/185 [=====] - 2s 9ms/step - loss: 1.0558 - accuracy: 0.6233 -  
val_loss: 1.0118 - val_accuracy: 0.6356  
Epoch 11/50  
185/185 [=====] - 1s 8ms/step - loss: 1.0144 - accuracy: 0.6387 -  
val_loss: 1.0045 - val_accuracy: 0.6471  
Epoch 12/50  
185/185 [=====] - 2s 8ms/step - loss: 0.9750 - accuracy: 0.6475 -  
val_loss: 0.9870 - val_accuracy: 0.6577  
Epoch 13/50  
185/185 [=====] - 2s 9ms/step - loss: 0.9228 - accuracy: 0.6621 -  
val_loss: 0.9446 - val_accuracy: 0.6731  
Epoch 14/50  
185/185 [=====] - 2s 8ms/step - loss: 0.8844 - accuracy: 0.6770 -  
val_loss: 0.9657 - val_accuracy: 0.6596  
Epoch 15/50  
185/185 [=====] - 1s 8ms/step - loss: 0.8395 - accuracy: 0.6936 -  
val_loss: 1.0189 - val_accuracy: 0.6808  
Epoch 16/50  
185/185 [=====] - 1s 8ms/step - loss: 0.8224 - accuracy: 0.6989 -  
val_loss: 0.9518 - val_accuracy: 0.6923  
Epoch 17/50  
185/185 [=====] - 2s 8ms/step - loss: 0.7970 - accuracy: 0.7104 -  
val_loss: 0.9277 - val_accuracy: 0.7067  
Epoch 18/50  
185/185 [=====] - 2s 8ms/step - loss: 0.7493 - accuracy: 0.7221 -  
val_loss: 0.9234 - val_accuracy: 0.6952  
Epoch 19/50  
185/185 [=====] - 2s 8ms/step - loss: 0.7205 - accuracy: 0.7353 -  
val_loss: 0.9911 - val_accuracy: 0.6779
```

Epoch 20/50
185/185 [=====] - 2s 8ms/step - loss: 0.7019 - accuracy: 0.7435 -
val_loss: 0.9394 - val_accuracy: 0.7029
Epoch 21/50
185/185 [=====] - 1s 8ms/step - loss: 0.7022 - accuracy: 0.7408 -
val_loss: 0.9890 - val_accuracy: 0.7115
Epoch 22/50
185/185 [=====] - 1s 7ms/step - loss: 0.6658 - accuracy: 0.7596 -
val_loss: 0.9891 - val_accuracy: 0.7058
Epoch 23/50
185/185 [=====] - 1s 7ms/step - loss: 0.6438 - accuracy: 0.7691 -
val_loss: 1.1430 - val_accuracy: 0.7058
Epoch 24/50
185/185 [=====] - 1s 7ms/step - loss: 0.6350 - accuracy: 0.7655 -
val_loss: 0.9763 - val_accuracy: 0.7163
Epoch 25/50
185/185 [=====] - 1s 7ms/step - loss: 0.6123 - accuracy: 0.7772 -
val_loss: 1.0541 - val_accuracy: 0.7096
Epoch 26/50
185/185 [=====] - 1s 7ms/step - loss: 0.5946 - accuracy: 0.7874 -
val_loss: 1.0307 - val_accuracy: 0.7212
Epoch 27/50
185/185 [=====] - 1s 7ms/step - loss: 0.5669 - accuracy: 0.7945 -
val_loss: 1.0550 - val_accuracy: 0.7192
Epoch 28/50
185/185 [=====] - 1s 7ms/step - loss: 0.5758 - accuracy: 0.7909 -
val_loss: 1.0178 - val_accuracy: 0.7308
Epoch 29/50
185/185 [=====] - 1s 7ms/step - loss: 0.5351 - accuracy: 0.8081 -
val_loss: 1.0167 - val_accuracy: 0.7260
Epoch 30/50
185/185 [=====] - 1s 7ms/step - loss: 0.5327 - accuracy: 0.8126 -
val_loss: 1.0278 - val_accuracy: 0.7192
Epoch 31/50
185/185 [=====] - 1s 7ms/step - loss: 0.5112 - accuracy: 0.8126 -
val_loss: 1.0080 - val_accuracy: 0.7317
Epoch 32/50
185/185 [=====] - 1s 8ms/step - loss: 0.5160 - accuracy: 0.8188 -
val_loss: 1.0006 - val_accuracy: 0.7337
Epoch 33/50
185/185 [=====] - 1s 7ms/step - loss: 0.4770 - accuracy: 0.8260 -
val_loss: 0.9891 - val_accuracy: 0.7433
Epoch 34/50
185/185 [=====] - 1s 7ms/step - loss: 0.4759 - accuracy: 0.8271 -
val_loss: 1.1083 - val_accuracy: 0.7510
Epoch 35/50
185/185 [=====] - 1s 7ms/step - loss: 0.4937 - accuracy: 0.8227 -
val_loss: 1.0128 - val_accuracy: 0.7346
Epoch 36/50
185/185 [=====] - 1s 7ms/step - loss: 0.4736 - accuracy: 0.8301 -
val_loss: 1.0744 - val_accuracy: 0.7394
Epoch 37/50
185/185 [=====] - 2s 10ms/step - loss: 0.4949 - accuracy: 0.8223 -
val_loss: 0.9941 - val_accuracy: 0.7337
Epoch 38/50
185/185 [=====] - 2s 8ms/step - loss: 0.4626 - accuracy: 0.8332 -
val_loss: 1.0812 - val_accuracy: 0.7375
Epoch 39/50
185/185 [=====] - 1s 7ms/step - loss: 0.4406 - accuracy: 0.8367 -
val_loss: 1.0585 - val_accuracy: 0.7413
Epoch 40/50
185/185 [=====] - 1s 7ms/step - loss: 0.4426 - accuracy: 0.8391 -
val_loss: 1.1254 - val_accuracy: 0.7375
Epoch 41/50
185/185 [=====] - 1s 8ms/step - loss: 0.4356 - accuracy: 0.8437 -
val_loss: 1.1296 - val_accuracy: 0.7490
Epoch 42/50
185/185 [=====] - 1s 7ms/step - loss: 0.4337 - accuracy: 0.8452 -
val_loss: 1.1849 - val_accuracy: 0.7356
Epoch 43/50
185/185 [=====] - 1s 7ms/step - loss: 0.3980 - accuracy: 0.8576 -
val_loss: 1.1357 - val_accuracy: 0.7269
Epoch 44/50
185/185 [=====] - 1s 7ms/step - loss: 0.3998 - accuracy: 0.8547 -

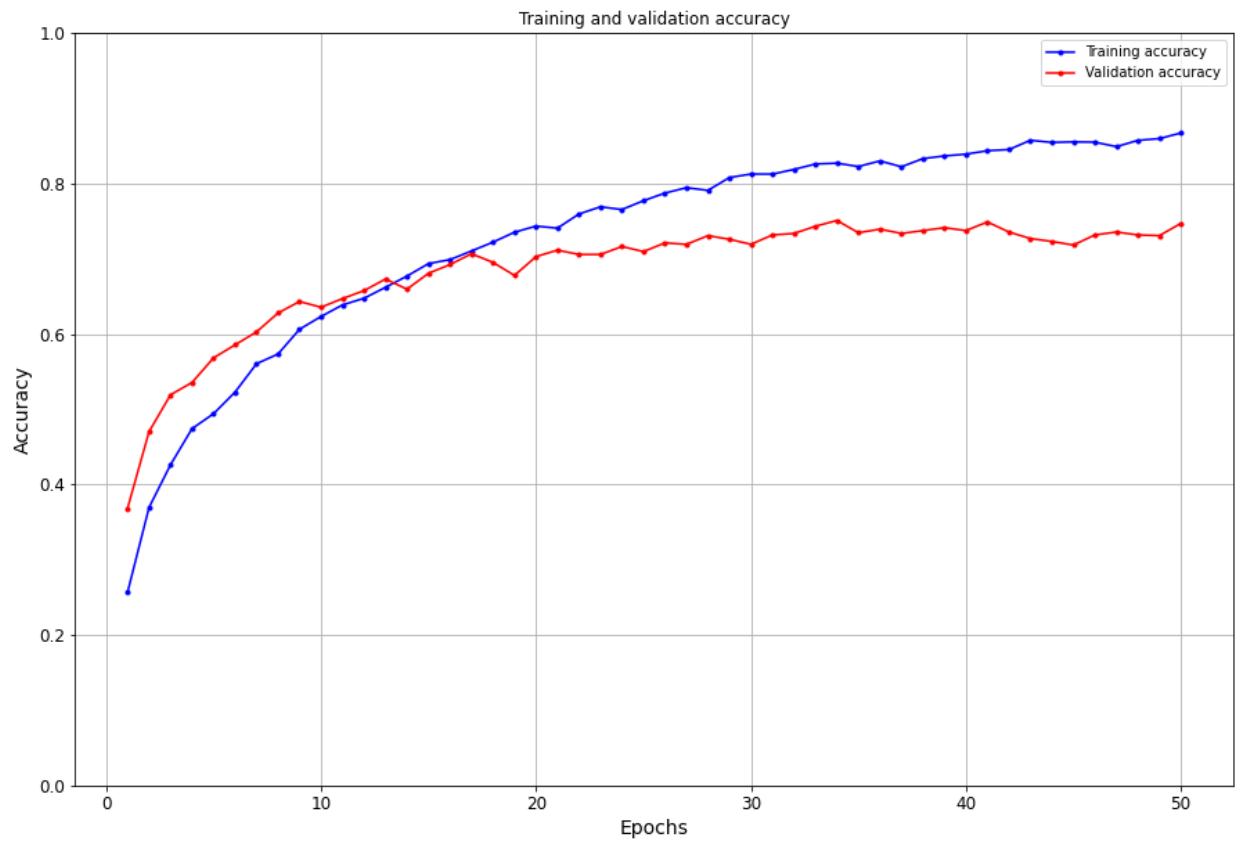
```
val_loss: 1.1806 - val_accuracy: 0.7231
Epoch 45/50
185/185 [=====] - 1s 7ms/step - loss: 0.4130 - accuracy: 0.8554 -
val_loss: 1.1374 - val_accuracy: 0.7183
Epoch 46/50
185/185 [=====] - 1s 7ms/step - loss: 0.4140 - accuracy: 0.8550 -
val_loss: 1.2345 - val_accuracy: 0.7317
Epoch 47/50
185/185 [=====] - 1s 8ms/step - loss: 0.4186 - accuracy: 0.8491 -
val_loss: 1.2398 - val_accuracy: 0.7356
Epoch 48/50
185/185 [=====] - 2s 9ms/step - loss: 0.4011 - accuracy: 0.8574 -
val_loss: 1.1644 - val_accuracy: 0.7317
Epoch 49/50
185/185 [=====] - 2s 9ms/step - loss: 0.3941 - accuracy: 0.8598 -
val_loss: 1.1217 - val_accuracy: 0.7308
Epoch 50/50
185/185 [=====] - 1s 7ms/step - loss: 0.3722 - accuracy: 0.8672 -
val_loss: 1.2449 - val_accuracy: 0.7471
```

```
In [ ]: plot_loss_and_accuracy(history)
```

```
accuracy    0.867243
dtype: float64 val_accuracy    0.750962
dtype: float64
```



```
<Figure size 432x288 with 0 Axes>
```



```
In [ ]: model.evaluate(test_ds)
```

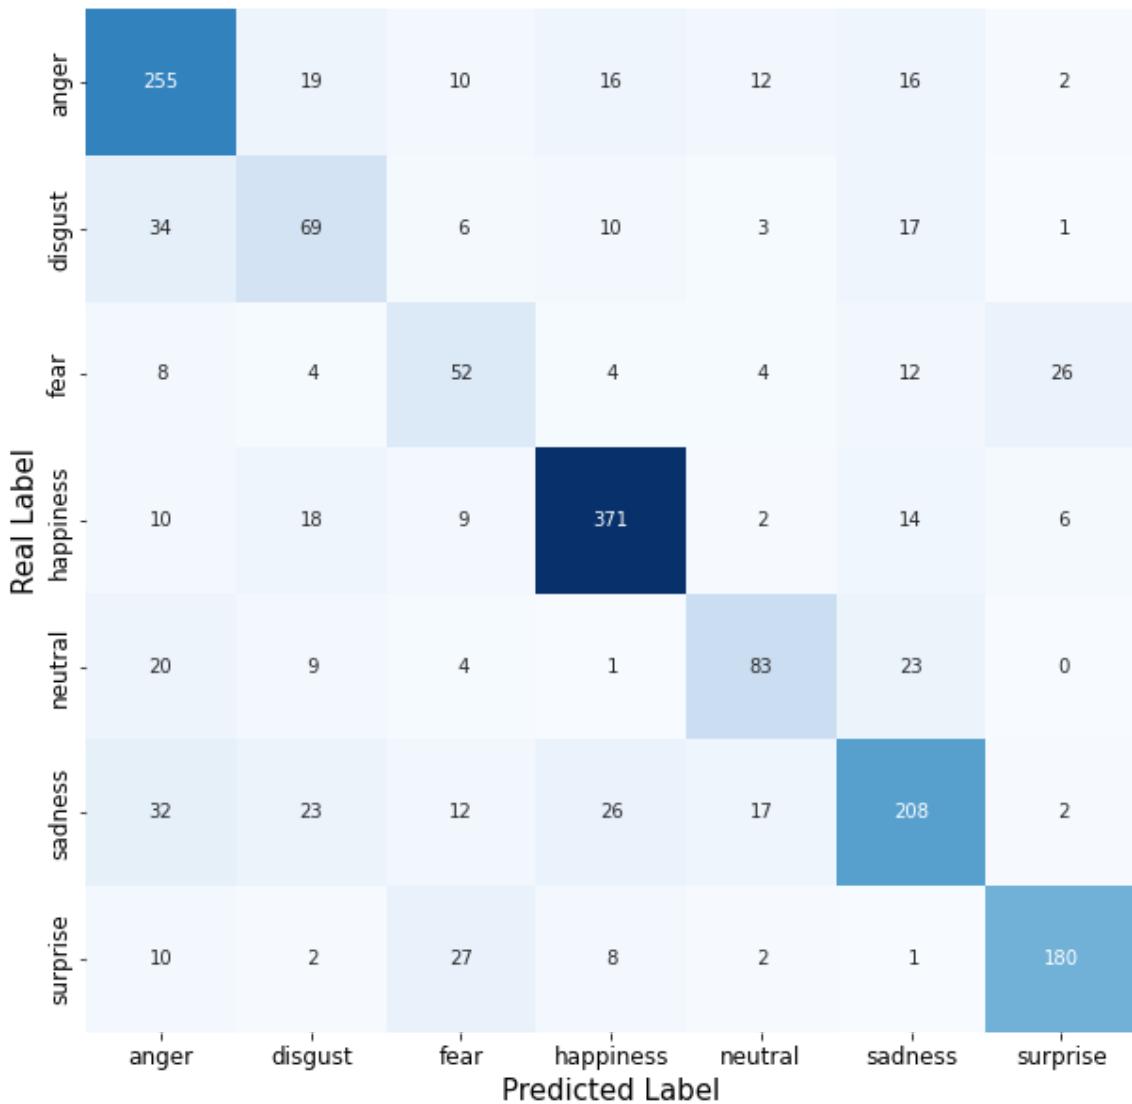
```
54/54 [=====] - 0s 4ms/step - loss: 1.5071 - accuracy: 0.7165
```

```
Out[ ]: [1.507069706916809, 0.7164705991744995]
```

```
In [ ]: # y_pred are probs, we have to convert them to int  
y_pred = model.predict(test_ds)
```

```
In [ ]: cnf_matrix = plot_confusion_matrix(y_pred, y_test, class_names, return_m=True)
```

Confusion Matrix



```
In [ ]: print_confusion_matrix_metrics(cnf_matrix)
```

False Positives:

```
anger -> 114
disgust -> 75
fear -> 68
happiness -> 65
neutral -> 40
sadness -> 83
surprise -> 37
```

False Negatives:

```
anger -> 75
disgust -> 71
fear -> 58
happiness -> 59
neutral -> 57
sadness -> 112
surprise -> 50
```

True Positives:

```
anger -> 255
disgust -> 69
fear -> 52
happiness -> 371
neutral -> 83
sadness -> 208
```

surprise -> 180

True Negatives:

anger -> 1256
disgust -> 1485
fear -> 1522
happiness -> 1205
neutral -> 1520
sadness -> 1297
surprise -> 1433

True Positive Rate:

anger -> 0.773%
disgust -> 0.493%
fear -> 0.473%
happiness -> 0.863%
neutral -> 0.593%
sadness -> 0.650%
surprise -> 0.783%

True Negative Rate:

anger -> 0.917%
disgust -> 0.952%
fear -> 0.957%
happiness -> 0.949%
neutral -> 0.974%
sadness -> 0.940%
surprise -> 0.975%

Positive Predictive Value:

anger -> 1
disgust -> 0
fear -> 0
happiness -> 1
neutral -> 1
sadness -> 1
surprise -> 1

Negative Predictive Value:

anger -> 1
disgust -> 1
fear -> 1
happiness -> 1
neutral -> 1
sadness -> 1
surprise -> 1

False Positive Rate:

anger -> 0.083%
disgust -> 0.048%
fear -> 0.043%
happiness -> 0.051%
neutral -> 0.026%
sadness -> 0.060%
surprise -> 0.025%

False Negative Rate:

anger -> 0.227%
disgust -> 0.507%
fear -> 0.527%
happiness -> 0.137%
neutral -> 0.407%
sadness -> 0.350%
surprise -> 0.217%

False Discovery Rate:

```
anger -> 0.309%
disgust -> 0.521%
fear -> 0.567%
happiness -> 0.149%
neutral -> 0.325%
sadness -> 0.285%
surprise -> 0.171%
```

Accuracy:

```
anger -> 0.889%
disgust -> 0.914%
fear -> 0.926%
happiness -> 0.927%
neutral -> 0.943%
sadness -> 0.885%
surprise -> 0.949%
```

```
In [ ]: print_general_metrics(y_pred, y_test)
```

Classification Report:

	precision	recall	f1-score	support
0	0.69	0.77	0.73	330
1	0.48	0.49	0.49	140
2	0.43	0.47	0.45	110
3	0.85	0.86	0.86	430
4	0.67	0.59	0.63	140
5	0.71	0.65	0.68	320
6	0.83	0.78	0.81	230
accuracy			0.72	1700
macro avg	0.67	0.66	0.66	1700
weighted avg	0.72	0.72	0.72	1700

Accuracy: 0.716

Balanced accuracy: 0.661

F1_Score: 0.717

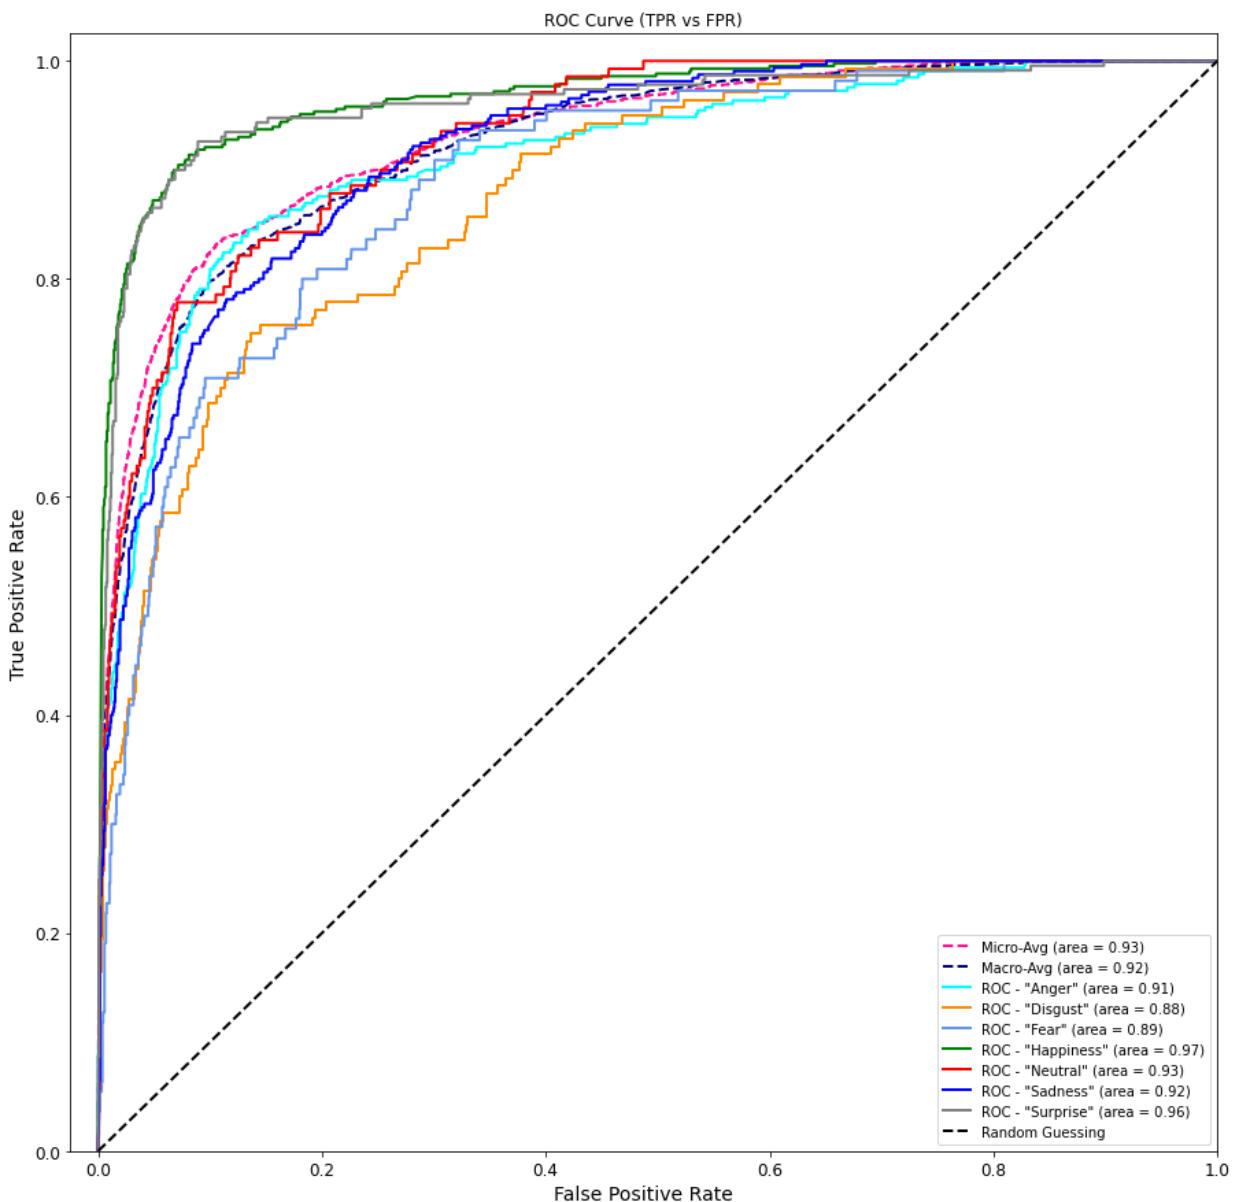
Precision: 0.719

Recall: 0.716

ROC AUC (OVR): 0.923

ROC AUC (OVO): 0.910

```
In [ ]: plot_multiclass_roc_curve(y_test, y_pred, class_names)
```



One-vs-One ROC AUC scores:

0.922703 (macro),
0.931220 (weighted by prevalence)

One-vs-Rest ROC AUC scores:

0.922703 (macro),
0.931220 (weighted by prevalence)

Create data augmentation

In []:

```
import random

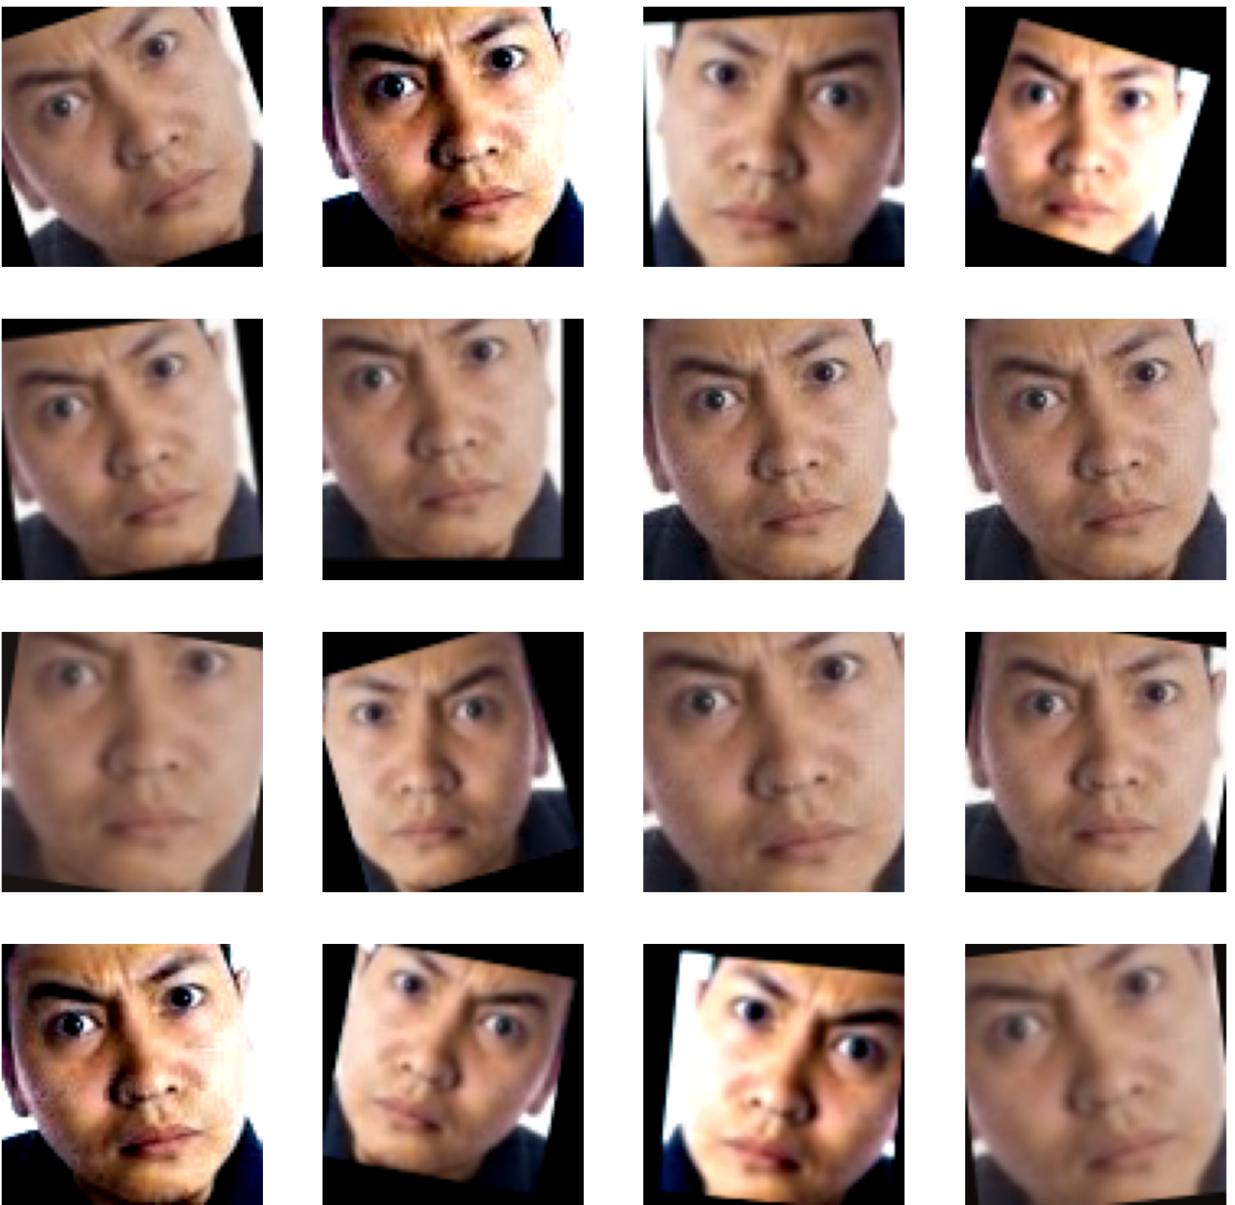
transformations = [
    layers.RandomFlip("horizontal", input_shape=(80, 80, 3)),
    layers.RandomTranslation(height_factor=0.1, width_factor=0.1, fill_mode="constant", input_shape=(80, 80, 3)),
    layers.RandomZoom(0.2, fill_mode="constant", input_shape=(80, 80, 3)),
    layers.RandomRotation(0.05, fill_mode="constant", input_shape=(80, 80, 3)),
    layers.RandomContrast(factor=[0.5, 1.3], input_shape=(80, 80, 3)),
    layers.GaussianNoise(0.03)
]

def customDataAugmentationLayer(input):
    if random.random() > 0.5 :
        return random.choice(transformations)(input)
    else:
        result = input
        for transformation in transformations:
            result = transformation(result)
```

```
return result

= layers.Lambda(lambda input: customDataAugmentationLayer(input))
```

```
In [ ]: plt.figure(figsize=(20, 20))
for images, _ in train_ds.take(1):
    for i in range(16):
        augmented_images = dataAug(images)
        #augmented_images = CustomDataAugmentationLayer()(images)
        ax = plt.subplot(4, 4, i + 1)
        plt.imshow(augmented_images[12].numpy().astype("float32"))
        plt.axis("off")
```



Improved CNN

In []:

```
keras.backend.clear_session()
tf.random.set_seed(42)
np.random.seed(42)

inputs = keras.Input(shape=(80, 80, 3))

x = keras.Sequential(transformations)(inputs)

x = layers.Conv2D(filters=32, kernel_size=3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.Conv2D(filters=32, kernel_size=3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.Conv2D(filters=32, kernel_size=3, padding='same', use_bias=False)(x)
x = layers.Activation("relu")(x)
x = layers.BatchNormalization()(x)
x = layers.MaxPooling2D()(x)

x = layers.Conv2D(64, 3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.Conv2D(64, 3, padding='same', use_bias=False)(x)
```

```

x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.Conv2D(64, 3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.MaxPooling2D()(x)

x = layers.Conv2D(128, 3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.Conv2D(128, 3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.Conv2D(128, 3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.MaxPooling2D()(x)

x = layers.Conv2D(256, 3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.Conv2D(256, 3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.Conv2D(256, 3, padding='same', use_bias=False)(x)
x = layers.BatchNormalization()(x)
x = layers.Activation("relu")(x)
x = layers.Dropout(0.1)(x)

x = layers.Flatten()(x)
x = layers.BatchNormalization()(x)
x = layers.Dense(128, activation='relu')(x)
x = layers.Dropout(0.25)(x)
x = layers.BatchNormalization()(x)
x = layers.Dense(60, activation='relu')(x)
x = layers.Dropout(0.25)(x)
x = layers.BatchNormalization()(x)
x = layers.Dense(30, activation='relu')(x)
x = layers.Dropout(0.25)(x)
x = layers.BatchNormalization()(x)

outputs = keras.layers.Dense(7, activation="softmax")(x)

modelWithAugmentation = keras.Model(inputs=inputs, outputs=outputs)
modelWithAugmentation.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	[(None, 80, 80, 3)]	0
sequential (Sequential)	(None, 80, 80, 3)	0
conv2d (Conv2D)	(None, 80, 80, 32)	864
batch_normalization (BatchNo	(None, 80, 80, 32)	128
activation (Activation)	(None, 80, 80, 32)	0
conv2d_1 (Conv2D)	(None, 80, 80, 32)	9216
batch_normalization_1 (Batch	(None, 80, 80, 32)	128
activation_1 (Activation)	(None, 80, 80, 32)	0

conv2d_2 (Conv2D)	(None, 80, 80, 32)	9216
activation_2 (Activation)	(None, 80, 80, 32)	0
batch_normalization_2 (Batch Normalization)	(None, 80, 80, 32)	128
max_pooling2d (MaxPooling2D)	(None, 40, 40, 32)	0
conv2d_3 (Conv2D)	(None, 40, 40, 64)	18432
batch_normalization_3 (Batch Normalization)	(None, 40, 40, 64)	256
activation_3 (Activation)	(None, 40, 40, 64)	0
conv2d_4 (Conv2D)	(None, 40, 40, 64)	36864
batch_normalization_4 (Batch Normalization)	(None, 40, 40, 64)	256
activation_4 (Activation)	(None, 40, 40, 64)	0
conv2d_5 (Conv2D)	(None, 40, 40, 64)	36864
batch_normalization_5 (Batch Normalization)	(None, 40, 40, 64)	256
activation_5 (Activation)	(None, 40, 40, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 20, 20, 64)	0
conv2d_6 (Conv2D)	(None, 20, 20, 128)	73728
batch_normalization_6 (Batch Normalization)	(None, 20, 20, 128)	512
activation_6 (Activation)	(None, 20, 20, 128)	0
conv2d_7 (Conv2D)	(None, 20, 20, 128)	147456
batch_normalization_7 (Batch Normalization)	(None, 20, 20, 128)	512
activation_7 (Activation)	(None, 20, 20, 128)	0
conv2d_8 (Conv2D)	(None, 20, 20, 128)	147456
batch_normalization_8 (Batch Normalization)	(None, 20, 20, 128)	512
activation_8 (Activation)	(None, 20, 20, 128)	0
max_pooling2d_2 (MaxPooling2D)	(None, 10, 10, 128)	0
conv2d_9 (Conv2D)	(None, 10, 10, 256)	294912
batch_normalization_9 (Batch Normalization)	(None, 10, 10, 256)	1024
activation_9 (Activation)	(None, 10, 10, 256)	0
conv2d_10 (Conv2D)	(None, 10, 10, 256)	589824
batch_normalization_10 (Batch Normalization)	(None, 10, 10, 256)	1024
activation_10 (Activation)	(None, 10, 10, 256)	0
conv2d_11 (Conv2D)	(None, 10, 10, 256)	589824
batch_normalization_11 (Batch Normalization)	(None, 10, 10, 256)	1024
activation_11 (Activation)	(None, 10, 10, 256)	0
dropout (Dropout)	(None, 10, 10, 256)	0
flatten (Flatten)	(None, 25600)	0
batch_normalization_12 (Batch Normalization)	(None, 25600)	102400
dense (Dense)	(None, 128)	3276928

dropout_1 (Dropout)	(None, 128)	0
batch_normalization_13 (BatchNormalization)	(None, 128)	512
dense_1 (Dense)	(None, 60)	7740
dropout_2 (Dropout)	(None, 60)	0
batch_normalization_14 (BatchNormalization)	(None, 60)	240
dense_2 (Dense)	(None, 30)	1830
dropout_3 (Dropout)	(None, 30)	0
batch_normalization_15 (BatchNormalization)	(None, 30)	120
dense_3 (Dense)	(None, 7)	217
<hr/>		
Total params: 5,350,403		
Trainable params: 5,295,887		
Non-trainable params: 54,516		

In []:

```
# Tensorboard Logs
root_logdir = os.path.join(os.curdir, 'tensorboard_logs_CNN')

# Sub-diretoria onde ficarão os dados da execução: tem um nome criado a partir da hora de inicio
def get_run_logdir():
    run_id = time.strftime("run_%Y_%m_%d-%H_%M_%S")
    return os.path.join(root_logdir, run_id)

best_model_file_path = "model_CNN.hdf5"

callbacks_list = [
    keras.callbacks.TensorBoard(
        get_run_logdir(),
        histogram_freq = 1
    ),
    tf.keras.callbacks.EarlyStopping(
        monitor='val_accuracy',
        patience = 500,
        min_delta = 0.005,
        restore_best_weights = True,
        mode='auto'
    ),
    keras.callbacks.ModelCheckpoint(
        filepath = best_model_file_path,
        monitor = "val_accuracy",
        verbose=0,
        save_best_only = True,
        save_weights_only = False,
        mode='max'
    )
]

# Compilar e treinar
#from keras.callbacks import ModelCheckpoint
#filepath = 'my_best_model.hdf5'
#checkpoint = ModelCheckpoint(filepath=filepath,
#                            monitor='val_accuracy',
#                            verbose=0,
#                            save_best_only=True,
#                            mode='max')
#callbacks = [checkpoint]
```

In []:

```
L = keras.losses.SparseCategoricalCrossentropy()

modelWithAugmentation.compile(loss=L, optimizer="adam", metrics=["accuracy"])
```

```
history = modelWithAugmentation.fit(  
    train_ds,  
    validation_data=val_ds,  
    epochs=200,  
    callbacks=callbacks_list  
)
```

```
Epoch 1/200  
185/185 [=====] - 10s 43ms/step - loss: 2.1864 - accuracy: 0.1931  
- val_loss: 3.4879 - val_accuracy: 0.1750  
Epoch 2/200  
185/185 [=====] - 7s 39ms/step - loss: 1.9369 - accuracy: 0.2599 -  
val_loss: 1.9714 - val_accuracy: 0.2500  
Epoch 3/200  
185/185 [=====] - 7s 39ms/step - loss: 1.7552 - accuracy: 0.3410 -  
val_loss: 1.8187 - val_accuracy: 0.3250  
Epoch 4/200  
185/185 [=====] - 7s 38ms/step - loss: 1.5977 - accuracy: 0.4137 -  
val_loss: 1.4470 - val_accuracy: 0.4894  
Epoch 5/200  
185/185 [=====] - 7s 38ms/step - loss: 1.4905 - accuracy: 0.4556 -  
val_loss: 1.4694 - val_accuracy: 0.4606  
Epoch 6/200  
185/185 [=====] - 7s 38ms/step - loss: 1.3889 - accuracy: 0.5015 -  
val_loss: 1.2969 - val_accuracy: 0.5356  
Epoch 7/200  
185/185 [=====] - 7s 38ms/step - loss: 1.3293 - accuracy: 0.5273 -  
val_loss: 1.7715 - val_accuracy: 0.3712  
Epoch 8/200  
185/185 [=====] - 7s 39ms/step - loss: 1.2618 - accuracy: 0.5544 -  
val_loss: 1.0919 - val_accuracy: 0.5990  
Epoch 9/200  
185/185 [=====] - 7s 38ms/step - loss: 1.2051 - accuracy: 0.5746 -  
val_loss: 1.6163 - val_accuracy: 0.4173  
Epoch 10/200  
185/185 [=====] - 7s 40ms/step - loss: 1.1541 - accuracy: 0.5899 -  
val_loss: 1.0474 - val_accuracy: 0.6173  
Epoch 11/200  
185/185 [=====] - 8s 42ms/step - loss: 1.1155 - accuracy: 0.6060 -  
val_loss: 1.0406 - val_accuracy: 0.5942  
Epoch 12/200  
185/185 [=====] - 7s 40ms/step - loss: 1.0851 - accuracy: 0.6121 -  
val_loss: 1.2565 - val_accuracy: 0.5375  
Epoch 13/200  
185/185 [=====] - 7s 39ms/step - loss: 1.0480 - accuracy: 0.6333 -  
val_loss: 1.0543 - val_accuracy: 0.5990  
Epoch 14/200  
185/185 [=====] - 8s 41ms/step - loss: 1.0353 - accuracy: 0.6392 -  
val_loss: 0.8821 - val_accuracy: 0.6933  
Epoch 15/200  
185/185 [=====] - 7s 40ms/step - loss: 0.9967 - accuracy: 0.6451 -  
val_loss: 0.9958 - val_accuracy: 0.6510  
Epoch 16/200  
185/185 [=====] - 7s 39ms/step - loss: 0.9881 - accuracy: 0.6553 -  
val_loss: 1.0804 - val_accuracy: 0.6010  
Epoch 17/200  
185/185 [=====] - 7s 39ms/step - loss: 0.9524 - accuracy: 0.6640 -  
val_loss: 0.9013 - val_accuracy: 0.6779  
Epoch 18/200  
185/185 [=====] - 8s 41ms/step - loss: 0.9482 - accuracy: 0.6628 -  
val_loss: 0.8775 - val_accuracy: 0.6846  
Epoch 19/200  
185/185 [=====] - 7s 40ms/step - loss: 0.9239 - accuracy: 0.6745 -  
val_loss: 1.0454 - val_accuracy: 0.6250  
Epoch 20/200  
185/185 [=====] - 7s 40ms/step - loss: 0.9055 - accuracy: 0.6765 -  
val_loss: 0.8936 - val_accuracy: 0.6740  
Epoch 21/200  
185/185 [=====] - 7s 40ms/step - loss: 0.8839 - accuracy: 0.6985 -  
val_loss: 0.8053 - val_accuracy: 0.7019  
Epoch 22/200  
185/185 [=====] - 7s 40ms/step - loss: 0.8884 - accuracy: 0.6909 -
```

```
val_loss: 0.7536 - val_accuracy: 0.7308
Epoch 23/200
185/185 [=====] - 7s 40ms/step - loss: 0.8724 - accuracy: 0.6980 -
val_loss: 0.8400 - val_accuracy: 0.6952
Epoch 24/200
185/185 [=====] - 7s 40ms/step - loss: 0.8669 - accuracy: 0.6962 -
val_loss: 0.7616 - val_accuracy: 0.7375
Epoch 25/200
185/185 [=====] - 8s 41ms/step - loss: 0.8321 - accuracy: 0.7062 -
val_loss: 0.7725 - val_accuracy: 0.7144
Epoch 26/200
185/185 [=====] - 7s 39ms/step - loss: 0.8431 - accuracy: 0.7121 -
val_loss: 0.8081 - val_accuracy: 0.7154
Epoch 27/200
185/185 [=====] - 7s 38ms/step - loss: 0.8129 - accuracy: 0.7163 -
val_loss: 1.0164 - val_accuracy: 0.6317
Epoch 28/200
185/185 [=====] - 7s 38ms/step - loss: 0.8261 - accuracy: 0.7175 -
val_loss: 0.7511 - val_accuracy: 0.7221
Epoch 29/200
185/185 [=====] - 7s 39ms/step - loss: 0.8100 - accuracy: 0.7230 -
val_loss: 1.0620 - val_accuracy: 0.6144
Epoch 30/200
185/185 [=====] - 7s 39ms/step - loss: 0.7800 - accuracy: 0.7335 -
val_loss: 1.0031 - val_accuracy: 0.6500
Epoch 31/200
185/185 [=====] - 7s 40ms/step - loss: 0.7753 - accuracy: 0.7333 -
val_loss: 0.8805 - val_accuracy: 0.6673
Epoch 32/200
185/185 [=====] - 8s 41ms/step - loss: 0.7727 - accuracy: 0.7306 -
val_loss: 0.7614 - val_accuracy: 0.7240
Epoch 33/200
185/185 [=====] - 7s 38ms/step - loss: 0.7388 - accuracy: 0.7460 -
val_loss: 0.8341 - val_accuracy: 0.6875
Epoch 34/200
185/185 [=====] - 8s 41ms/step - loss: 0.7578 - accuracy: 0.7391 -
val_loss: 0.7842 - val_accuracy: 0.7231
Epoch 35/200
185/185 [=====] - 7s 39ms/step - loss: 0.7235 - accuracy: 0.7525 -
val_loss: 0.7893 - val_accuracy: 0.7317
Epoch 36/200
185/185 [=====] - 7s 39ms/step - loss: 0.7247 - accuracy: 0.7528 -
val_loss: 0.6822 - val_accuracy: 0.7567
Epoch 37/200
185/185 [=====] - 7s 40ms/step - loss: 0.7219 - accuracy: 0.7553 -
val_loss: 0.7777 - val_accuracy: 0.7375
Epoch 38/200
185/185 [=====] - 7s 40ms/step - loss: 0.7044 - accuracy: 0.7628 -
val_loss: 0.7745 - val_accuracy: 0.7231
Epoch 39/200
185/185 [=====] - 7s 39ms/step - loss: 0.6888 - accuracy: 0.7603 -
val_loss: 0.7153 - val_accuracy: 0.7548
Epoch 40/200
185/185 [=====] - 7s 39ms/step - loss: 0.6962 - accuracy: 0.7679 -
val_loss: 0.8409 - val_accuracy: 0.7154
Epoch 41/200
185/185 [=====] - 7s 38ms/step - loss: 0.7018 - accuracy: 0.7648 -
val_loss: 0.9268 - val_accuracy: 0.6769
Epoch 42/200
185/185 [=====] - 7s 38ms/step - loss: 0.6745 - accuracy: 0.7791 -
val_loss: 0.6642 - val_accuracy: 0.7596
Epoch 43/200
185/185 [=====] - 7s 38ms/step - loss: 0.6566 - accuracy: 0.7803 -
val_loss: 0.6709 - val_accuracy: 0.7808
Epoch 44/200
185/185 [=====] - 7s 38ms/step - loss: 0.6705 - accuracy: 0.7786 -
val_loss: 0.7620 - val_accuracy: 0.7260
Epoch 45/200
185/185 [=====] - 7s 40ms/step - loss: 0.6438 - accuracy: 0.7818 -
val_loss: 0.6918 - val_accuracy: 0.7606
Epoch 46/200
185/185 [=====] - 7s 39ms/step - loss: 0.6354 - accuracy: 0.7887 -
val_loss: 0.6468 - val_accuracy: 0.7817
Epoch 47/200
```

185/185 [=====] - 7s 39ms/step - loss: 0.6177 - accuracy: 0.7913 -
val_loss: 0.6483 - val_accuracy: 0.7846
Epoch 48/200
185/185 [=====] - 7s 38ms/step - loss: 0.6277 - accuracy: 0.7928 -
val_loss: 0.6563 - val_accuracy: 0.7933
Epoch 49/200
185/185 [=====] - 8s 43ms/step - loss: 0.6106 - accuracy: 0.7967 -
val_loss: 0.7232 - val_accuracy: 0.7779
Epoch 50/200
185/185 [=====] - 8s 42ms/step - loss: 0.6095 - accuracy: 0.7984 -
val_loss: 0.8676 - val_accuracy: 0.7125
Epoch 51/200
185/185 [=====] - 8s 42ms/step - loss: 0.5954 - accuracy: 0.8052 -
val_loss: 0.6476 - val_accuracy: 0.7808
Epoch 52/200
185/185 [=====] - 8s 41ms/step - loss: 0.5910 - accuracy: 0.8071 -
val_loss: 0.6660 - val_accuracy: 0.7837
Epoch 53/200
185/185 [=====] - 7s 39ms/step - loss: 0.6089 - accuracy: 0.7998 -
val_loss: 0.6201 - val_accuracy: 0.8058
Epoch 54/200
185/185 [=====] - 7s 40ms/step - loss: 0.5857 - accuracy: 0.8018 -
val_loss: 0.6290 - val_accuracy: 0.7971
Epoch 55/200
185/185 [=====] - 7s 40ms/step - loss: 0.5642 - accuracy: 0.8142 -
val_loss: 0.7420 - val_accuracy: 0.7683
Epoch 56/200
185/185 [=====] - 7s 38ms/step - loss: 0.5629 - accuracy: 0.8165 -
val_loss: 0.6990 - val_accuracy: 0.7760
Epoch 57/200
185/185 [=====] - 7s 40ms/step - loss: 0.5592 - accuracy: 0.8154 -
val_loss: 0.6734 - val_accuracy: 0.7827
Epoch 58/200
185/185 [=====] - 7s 38ms/step - loss: 0.5444 - accuracy: 0.8216 -
val_loss: 0.5701 - val_accuracy: 0.8106
Epoch 59/200
185/185 [=====] - 7s 39ms/step - loss: 0.5290 - accuracy: 0.8277 -
val_loss: 0.6060 - val_accuracy: 0.8067
Epoch 60/200
185/185 [=====] - 8s 42ms/step - loss: 0.5440 - accuracy: 0.8186 -
val_loss: 0.7763 - val_accuracy: 0.7625
Epoch 61/200
185/185 [=====] - 7s 39ms/step - loss: 0.5280 - accuracy: 0.8301 -
val_loss: 0.6297 - val_accuracy: 0.7971
Epoch 62/200
185/185 [=====] - 8s 42ms/step - loss: 0.5177 - accuracy: 0.8366 -
val_loss: 0.6561 - val_accuracy: 0.7817
Epoch 63/200
185/185 [=====] - 7s 40ms/step - loss: 0.5241 - accuracy: 0.8330 -
val_loss: 0.6448 - val_accuracy: 0.7798
Epoch 64/200
185/185 [=====] - 7s 40ms/step - loss: 0.5192 - accuracy: 0.8316 -
val_loss: 0.6549 - val_accuracy: 0.7865
Epoch 65/200
185/185 [=====] - 7s 39ms/step - loss: 0.5067 - accuracy: 0.8372 -
val_loss: 0.6761 - val_accuracy: 0.7702
Epoch 66/200
185/185 [=====] - 7s 39ms/step - loss: 0.4861 - accuracy: 0.8440 -
val_loss: 0.6843 - val_accuracy: 0.7788
Epoch 67/200
185/185 [=====] - 7s 39ms/step - loss: 0.4819 - accuracy: 0.8438 -
val_loss: 0.6266 - val_accuracy: 0.7971
Epoch 68/200
185/185 [=====] - 7s 39ms/step - loss: 0.4865 - accuracy: 0.8452 -
val_loss: 0.5607 - val_accuracy: 0.8269
Epoch 69/200
185/185 [=====] - 7s 39ms/step - loss: 0.4716 - accuracy: 0.8511 -
val_loss: 0.6493 - val_accuracy: 0.7952
Epoch 70/200
185/185 [=====] - 7s 39ms/step - loss: 0.4612 - accuracy: 0.8499 -
val_loss: 0.6937 - val_accuracy: 0.7971
Epoch 71/200
185/185 [=====] - 8s 41ms/step - loss: 0.4684 - accuracy: 0.8496 -
val_loss: 0.7955 - val_accuracy: 0.7798

Epoch 72/200
185/185 [=====] - 7s 39ms/step - loss: 0.4631 - accuracy: 0.8508 -
val_loss: 0.7537 - val_accuracy: 0.7615
Epoch 73/200
185/185 [=====] - 7s 40ms/step - loss: 0.4725 - accuracy: 0.8523 -
val_loss: 0.7209 - val_accuracy: 0.7750
Epoch 74/200
185/185 [=====] - 7s 40ms/step - loss: 0.4476 - accuracy: 0.8594 -
val_loss: 0.6550 - val_accuracy: 0.7923
Epoch 75/200
185/185 [=====] - 8s 41ms/step - loss: 0.4470 - accuracy: 0.8571 -
val_loss: 0.7885 - val_accuracy: 0.7596
Epoch 76/200
185/185 [=====] - 7s 39ms/step - loss: 0.4253 - accuracy: 0.8700 -
val_loss: 0.8155 - val_accuracy: 0.7548
Epoch 77/200
185/185 [=====] - 7s 39ms/step - loss: 0.4431 - accuracy: 0.8618 -
val_loss: 0.6908 - val_accuracy: 0.7808
Epoch 78/200
185/185 [=====] - 7s 39ms/step - loss: 0.4179 - accuracy: 0.8681 -
val_loss: 0.6445 - val_accuracy: 0.7904
Epoch 79/200
185/185 [=====] - 7s 39ms/step - loss: 0.4261 - accuracy: 0.8616 -
val_loss: 0.9311 - val_accuracy: 0.7260
Epoch 80/200
185/185 [=====] - 7s 39ms/step - loss: 0.4317 - accuracy: 0.8605 -
val_loss: 0.6703 - val_accuracy: 0.7933
Epoch 81/200
185/185 [=====] - 7s 39ms/step - loss: 0.4072 - accuracy: 0.8696 -
val_loss: 0.7347 - val_accuracy: 0.7779
Epoch 82/200
185/185 [=====] - 7s 38ms/step - loss: 0.4035 - accuracy: 0.8728 -
val_loss: 0.7142 - val_accuracy: 0.7779
Epoch 83/200
185/185 [=====] - 7s 38ms/step - loss: 0.3868 - accuracy: 0.8835 -
val_loss: 0.7019 - val_accuracy: 0.7846
Epoch 84/200
185/185 [=====] - 7s 39ms/step - loss: 0.3964 - accuracy: 0.8693 -
val_loss: 0.6793 - val_accuracy: 0.8010
Epoch 85/200
185/185 [=====] - 7s 40ms/step - loss: 0.3888 - accuracy: 0.8784 -
val_loss: 0.6179 - val_accuracy: 0.8173
Epoch 86/200
185/185 [=====] - 7s 38ms/step - loss: 0.3965 - accuracy: 0.8713 -
val_loss: 0.6326 - val_accuracy: 0.8029
Epoch 87/200
185/185 [=====] - 7s 38ms/step - loss: 0.3764 - accuracy: 0.8776 -
val_loss: 0.6829 - val_accuracy: 0.7923
Epoch 88/200
185/185 [=====] - 7s 39ms/step - loss: 0.3870 - accuracy: 0.8795 -
val_loss: 0.5837 - val_accuracy: 0.8144
Epoch 89/200
185/185 [=====] - 7s 38ms/step - loss: 0.3902 - accuracy: 0.8801 -
val_loss: 0.5958 - val_accuracy: 0.8154
Epoch 90/200
185/185 [=====] - 7s 39ms/step - loss: 0.3683 - accuracy: 0.8840 -
val_loss: 0.6259 - val_accuracy: 0.8067
Epoch 91/200
185/185 [=====] - 7s 39ms/step - loss: 0.3575 - accuracy: 0.8861 -
val_loss: 0.6639 - val_accuracy: 0.7990
Epoch 92/200
185/185 [=====] - 7s 38ms/step - loss: 0.3702 - accuracy: 0.8756 -
val_loss: 0.6503 - val_accuracy: 0.8067
Epoch 93/200
185/185 [=====] - 7s 39ms/step - loss: 0.3558 - accuracy: 0.8849 -
val_loss: 0.6849 - val_accuracy: 0.7837
Epoch 94/200
185/185 [=====] - 8s 41ms/step - loss: 0.3611 - accuracy: 0.8864 -
val_loss: 0.6155 - val_accuracy: 0.8269
Epoch 95/200
185/185 [=====] - 7s 39ms/step - loss: 0.3638 - accuracy: 0.8889 -
val_loss: 0.5887 - val_accuracy: 0.8106
Epoch 96/200
185/185 [=====] - 7s 39ms/step - loss: 0.3256 - accuracy: 0.8971 -

```
val_loss: 0.6124 - val_accuracy: 0.8260
Epoch 97/200
185/185 [=====] - 7s 38ms/step - loss: 0.3408 - accuracy: 0.8893 -
val_loss: 0.6234 - val_accuracy: 0.8250
Epoch 98/200
185/185 [=====] - 7s 38ms/step - loss: 0.3340 - accuracy: 0.8874 -
val_loss: 0.6326 - val_accuracy: 0.8240
Epoch 99/200
185/185 [=====] - 7s 40ms/step - loss: 0.3506 - accuracy: 0.8900 -
val_loss: 0.5931 - val_accuracy: 0.8202
Epoch 100/200
185/185 [=====] - 8s 42ms/step - loss: 0.3485 - accuracy: 0.8861 -
val_loss: 0.6724 - val_accuracy: 0.8038
Epoch 101/200
185/185 [=====] - 7s 39ms/step - loss: 0.3325 - accuracy: 0.8967 -
val_loss: 0.6317 - val_accuracy: 0.8115
Epoch 102/200
185/185 [=====] - 7s 39ms/step - loss: 0.3296 - accuracy: 0.8940 -
val_loss: 0.6385 - val_accuracy: 0.8058
Epoch 103/200
185/185 [=====] - 7s 39ms/step - loss: 0.3752 - accuracy: 0.8822 -
val_loss: 0.5690 - val_accuracy: 0.8404
Epoch 104/200
185/185 [=====] - 7s 38ms/step - loss: 0.3127 - accuracy: 0.9086 -
val_loss: 0.6662 - val_accuracy: 0.8221
Epoch 105/200
185/185 [=====] - 7s 38ms/step - loss: 0.3188 - accuracy: 0.9001 -
val_loss: 0.6718 - val_accuracy: 0.8154
Epoch 106/200
185/185 [=====] - 7s 38ms/step - loss: 0.3158 - accuracy: 0.8993 -
val_loss: 0.6637 - val_accuracy: 0.8087
Epoch 107/200
185/185 [=====] - 7s 39ms/step - loss: 0.2993 - accuracy: 0.9067 -
val_loss: 0.8665 - val_accuracy: 0.7500
Epoch 108/200
185/185 [=====] - 7s 39ms/step - loss: 0.3087 - accuracy: 0.9040 -
val_loss: 0.6094 - val_accuracy: 0.8202
Epoch 109/200
185/185 [=====] - 7s 39ms/step - loss: 0.3010 - accuracy: 0.9062 -
val_loss: 0.5994 - val_accuracy: 0.8423
Epoch 110/200
185/185 [=====] - 7s 38ms/step - loss: 0.2969 - accuracy: 0.9066 -
val_loss: 0.6212 - val_accuracy: 0.8250
Epoch 111/200
185/185 [=====] - 7s 38ms/step - loss: 0.2981 - accuracy: 0.9057 -
val_loss: 0.5845 - val_accuracy: 0.8308
Epoch 112/200
185/185 [=====] - 7s 38ms/step - loss: 0.2799 - accuracy: 0.9076 -
val_loss: 0.7031 - val_accuracy: 0.8115
Epoch 113/200
185/185 [=====] - 7s 38ms/step - loss: 0.2912 - accuracy: 0.9093 -
val_loss: 0.5691 - val_accuracy: 0.8481
Epoch 114/200
185/185 [=====] - 7s 39ms/step - loss: 0.2789 - accuracy: 0.9157 -
val_loss: 0.6623 - val_accuracy: 0.8260
Epoch 115/200
185/185 [=====] - 7s 38ms/step - loss: 0.2925 - accuracy: 0.9095 -
val_loss: 0.6653 - val_accuracy: 0.8269
Epoch 116/200
185/185 [=====] - 7s 38ms/step - loss: 0.2851 - accuracy: 0.9140 -
val_loss: 0.7601 - val_accuracy: 0.7962
Epoch 117/200
185/185 [=====] - 7s 38ms/step - loss: 0.2653 - accuracy: 0.9169 -
val_loss: 0.5976 - val_accuracy: 0.8404
Epoch 118/200
185/185 [=====] - 7s 38ms/step - loss: 0.2701 - accuracy: 0.9208 -
val_loss: 0.6571 - val_accuracy: 0.8192
Epoch 119/200
185/185 [=====] - 7s 39ms/step - loss: 0.2780 - accuracy: 0.9122 -
val_loss: 0.7117 - val_accuracy: 0.8019
Epoch 120/200
185/185 [=====] - 7s 39ms/step - loss: 0.2715 - accuracy: 0.9129 -
val_loss: 0.6747 - val_accuracy: 0.8106
Epoch 121/200
```

185/185 [=====] - 7s 39ms/step - loss: 0.2720 - accuracy: 0.9157 -
val_loss: 0.5825 - val_accuracy: 0.8375
Epoch 122/200
185/185 [=====] - 7s 39ms/step - loss: 0.2434 - accuracy: 0.9237 -
val_loss: 0.6496 - val_accuracy: 0.8173
Epoch 123/200
185/185 [=====] - 7s 38ms/step - loss: 0.2529 - accuracy: 0.9195 -
val_loss: 0.7685 - val_accuracy: 0.8058
Epoch 124/200
185/185 [=====] - 7s 38ms/step - loss: 0.2800 - accuracy: 0.9130 -
val_loss: 0.7392 - val_accuracy: 0.8125
Epoch 125/200
185/185 [=====] - 7s 39ms/step - loss: 0.2584 - accuracy: 0.9168 -
val_loss: 0.7034 - val_accuracy: 0.8192
Epoch 126/200
185/185 [=====] - 7s 39ms/step - loss: 0.2601 - accuracy: 0.9164 -
val_loss: 0.7650 - val_accuracy: 0.7865
Epoch 127/200
185/185 [=====] - 7s 39ms/step - loss: 0.2626 - accuracy: 0.9156 -
val_loss: 0.6482 - val_accuracy: 0.8279
Epoch 128/200
185/185 [=====] - 7s 39ms/step - loss: 0.2530 - accuracy: 0.9223 -
val_loss: 0.7802 - val_accuracy: 0.7962
Epoch 129/200
185/185 [=====] - 7s 39ms/step - loss: 0.2755 - accuracy: 0.9147 -
val_loss: 0.7629 - val_accuracy: 0.8077
Epoch 130/200
185/185 [=====] - 7s 38ms/step - loss: 0.2462 - accuracy: 0.9249 -
val_loss: 0.7614 - val_accuracy: 0.8038
Epoch 131/200
185/185 [=====] - 7s 38ms/step - loss: 0.2380 - accuracy: 0.9242 -
val_loss: 0.7683 - val_accuracy: 0.8019
Epoch 132/200
185/185 [=====] - 7s 38ms/step - loss: 0.2252 - accuracy: 0.9315 -
val_loss: 0.6250 - val_accuracy: 0.8404
Epoch 133/200
185/185 [=====] - 7s 39ms/step - loss: 0.2450 - accuracy: 0.9210 -
val_loss: 0.7229 - val_accuracy: 0.8115
Epoch 134/200
185/185 [=====] - 7s 39ms/step - loss: 0.2307 - accuracy: 0.9271 -
val_loss: 0.7654 - val_accuracy: 0.8125
Epoch 135/200
185/185 [=====] - 7s 41ms/step - loss: 0.2320 - accuracy: 0.9235 -
val_loss: 0.8556 - val_accuracy: 0.7769
Epoch 136/200
185/185 [=====] - 7s 39ms/step - loss: 0.2405 - accuracy: 0.9213 -
val_loss: 0.6965 - val_accuracy: 0.8087
Epoch 137/200
185/185 [=====] - 7s 39ms/step - loss: 0.2277 - accuracy: 0.9323 -
val_loss: 0.7123 - val_accuracy: 0.8212
Epoch 138/200
185/185 [=====] - 7s 40ms/step - loss: 0.2350 - accuracy: 0.9283 -
val_loss: 0.7450 - val_accuracy: 0.8029
Epoch 139/200
185/185 [=====] - 7s 39ms/step - loss: 0.2154 - accuracy: 0.9352 -
val_loss: 0.7582 - val_accuracy: 0.7904
Epoch 140/200
185/185 [=====] - 7s 40ms/step - loss: 0.2338 - accuracy: 0.9317 -
val_loss: 0.6779 - val_accuracy: 0.8260
Epoch 141/200
185/185 [=====] - 7s 39ms/step - loss: 0.2004 - accuracy: 0.9344 -
val_loss: 0.6465 - val_accuracy: 0.8452
Epoch 142/200
185/185 [=====] - 7s 39ms/step - loss: 0.2043 - accuracy: 0.9349 -
val_loss: 0.6975 - val_accuracy: 0.8125
Epoch 143/200
185/185 [=====] - 7s 39ms/step - loss: 0.2065 - accuracy: 0.9351 -
val_loss: 0.6595 - val_accuracy: 0.8385
Epoch 144/200
185/185 [=====] - 7s 39ms/step - loss: 0.2095 - accuracy: 0.9351 -
val_loss: 0.8784 - val_accuracy: 0.7913
Epoch 145/200
185/185 [=====] - 7s 38ms/step - loss: 0.1927 - accuracy: 0.9417 -
val_loss: 0.8297 - val_accuracy: 0.8029

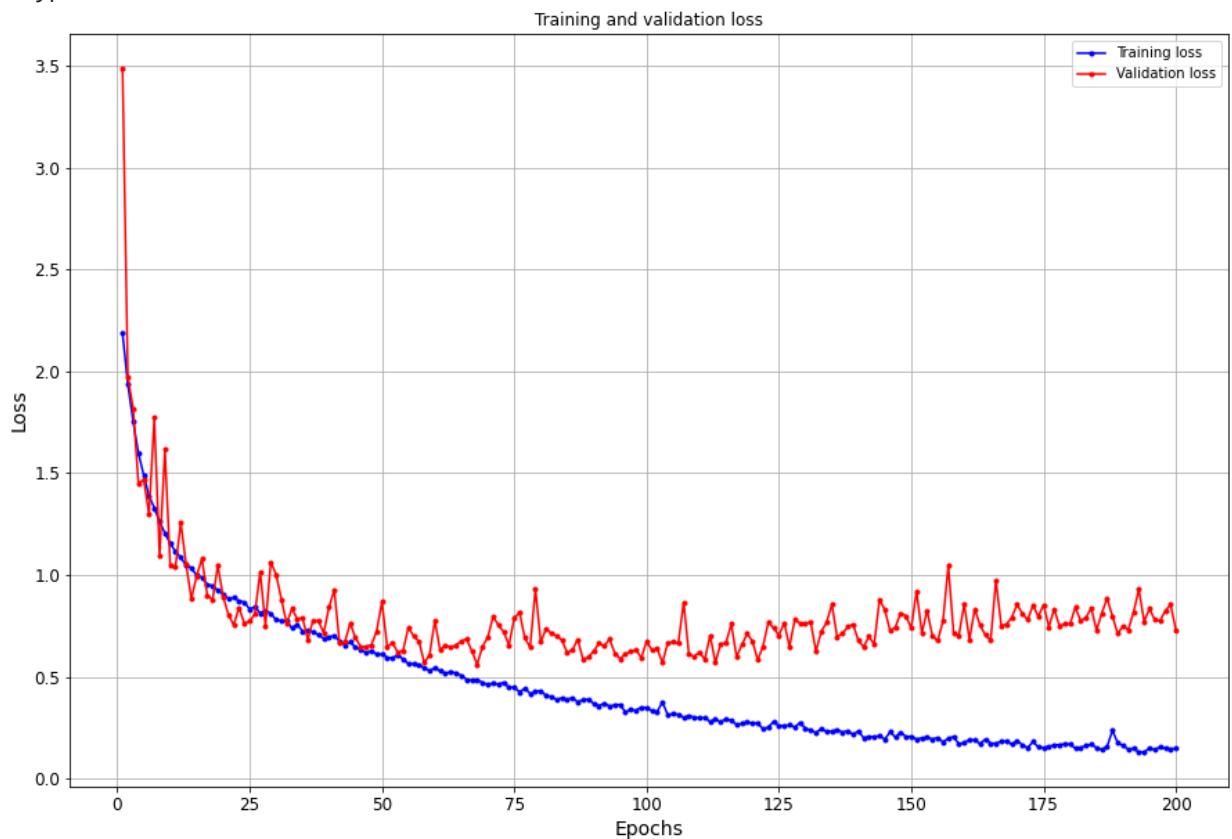
Epoch 146/200
185/185 [=====] - 7s 38ms/step - loss: 0.2291 - accuracy: 0.9291 -
val_loss: 0.7247 - val_accuracy: 0.8279
Epoch 147/200
185/185 [=====] - 7s 38ms/step - loss: 0.2030 - accuracy: 0.9386 -
val_loss: 0.7440 - val_accuracy: 0.8298
Epoch 148/200
185/185 [=====] - 7s 39ms/step - loss: 0.2282 - accuracy: 0.9276 -
val_loss: 0.8123 - val_accuracy: 0.8000
Epoch 149/200
185/185 [=====] - 7s 39ms/step - loss: 0.2054 - accuracy: 0.9359 -
val_loss: 0.7994 - val_accuracy: 0.8163
Epoch 150/200
185/185 [=====] - 7s 39ms/step - loss: 0.2081 - accuracy: 0.9362 -
val_loss: 0.7385 - val_accuracy: 0.8317
Epoch 151/200
185/185 [=====] - 7s 38ms/step - loss: 0.1935 - accuracy: 0.9408 -
val_loss: 0.9211 - val_accuracy: 0.8067
Epoch 152/200
185/185 [=====] - 7s 39ms/step - loss: 0.1990 - accuracy: 0.9410 -
val_loss: 0.7150 - val_accuracy: 0.8288
Epoch 153/200
185/185 [=====] - 7s 40ms/step - loss: 0.2020 - accuracy: 0.9330 -
val_loss: 0.8248 - val_accuracy: 0.8202
Epoch 154/200
185/185 [=====] - 8s 41ms/step - loss: 0.1940 - accuracy: 0.9390 -
val_loss: 0.7000 - val_accuracy: 0.8337
Epoch 155/200
185/185 [=====] - 8s 41ms/step - loss: 0.2008 - accuracy: 0.9405 -
val_loss: 0.6773 - val_accuracy: 0.8317
Epoch 156/200
185/185 [=====] - 8s 41ms/step - loss: 0.1793 - accuracy: 0.9442 -
val_loss: 0.7745 - val_accuracy: 0.8250
Epoch 157/200
185/185 [=====] - 7s 39ms/step - loss: 0.1979 - accuracy: 0.9356 -
val_loss: 1.0454 - val_accuracy: 0.7558
Epoch 158/200
185/185 [=====] - 7s 38ms/step - loss: 0.2061 - accuracy: 0.9371 -
val_loss: 0.7147 - val_accuracy: 0.8375
Epoch 159/200
185/185 [=====] - 7s 38ms/step - loss: 0.1690 - accuracy: 0.9498 -
val_loss: 0.7038 - val_accuracy: 0.8288
Epoch 160/200
185/185 [=====] - 7s 38ms/step - loss: 0.1767 - accuracy: 0.9461 -
val_loss: 0.8572 - val_accuracy: 0.8019
Epoch 161/200
185/185 [=====] - 7s 38ms/step - loss: 0.1905 - accuracy: 0.9422 -
val_loss: 0.6826 - val_accuracy: 0.8250
Epoch 162/200
185/185 [=====] - 7s 38ms/step - loss: 0.1883 - accuracy: 0.9413 -
val_loss: 0.8330 - val_accuracy: 0.8087
Epoch 163/200
185/185 [=====] - 7s 38ms/step - loss: 0.1717 - accuracy: 0.9486 -
val_loss: 0.7523 - val_accuracy: 0.8202
Epoch 164/200
185/185 [=====] - 7s 38ms/step - loss: 0.1935 - accuracy: 0.9422 -
val_loss: 0.7055 - val_accuracy: 0.8462
Epoch 165/200
185/185 [=====] - 7s 38ms/step - loss: 0.1726 - accuracy: 0.9468 -
val_loss: 0.6793 - val_accuracy: 0.8404
Epoch 166/200
185/185 [=====] - 7s 39ms/step - loss: 0.1740 - accuracy: 0.9452 -
val_loss: 0.9737 - val_accuracy: 0.7894
Epoch 167/200
185/185 [=====] - 7s 38ms/step - loss: 0.1819 - accuracy: 0.9422 -
val_loss: 0.7510 - val_accuracy: 0.8317
Epoch 168/200
185/185 [=====] - 7s 38ms/step - loss: 0.1826 - accuracy: 0.9418 -
val_loss: 0.7585 - val_accuracy: 0.8260
Epoch 169/200
185/185 [=====] - 7s 38ms/step - loss: 0.1684 - accuracy: 0.9449 -
val_loss: 0.7918 - val_accuracy: 0.8250
Epoch 170/200
185/185 [=====] - 7s 38ms/step - loss: 0.1857 - accuracy: 0.9408 -

```
val_loss: 0.8549 - val_accuracy: 0.8135
Epoch 171/200
185/185 [=====] - 7s 38ms/step - loss: 0.1652 - accuracy: 0.9517 -
val_loss: 0.8092 - val_accuracy: 0.8212
Epoch 172/200
185/185 [=====] - 7s 38ms/step - loss: 0.1516 - accuracy: 0.9544 -
val_loss: 0.7800 - val_accuracy: 0.8404
Epoch 173/200
185/185 [=====] - 7s 38ms/step - loss: 0.1811 - accuracy: 0.9440 -
val_loss: 0.8512 - val_accuracy: 0.8106
Epoch 174/200
185/185 [=====] - 7s 38ms/step - loss: 0.1545 - accuracy: 0.9535 -
val_loss: 0.7965 - val_accuracy: 0.7990
Epoch 175/200
185/185 [=====] - 7s 38ms/step - loss: 0.1512 - accuracy: 0.9515 -
val_loss: 0.8537 - val_accuracy: 0.8279
Epoch 176/200
185/185 [=====] - 7s 38ms/step - loss: 0.1585 - accuracy: 0.9522 -
val_loss: 0.7448 - val_accuracy: 0.8346
Epoch 177/200
185/185 [=====] - 7s 38ms/step - loss: 0.1638 - accuracy: 0.9496 -
val_loss: 0.8331 - val_accuracy: 0.8327
Epoch 178/200
185/185 [=====] - 7s 38ms/step - loss: 0.1664 - accuracy: 0.9486 -
val_loss: 0.7492 - val_accuracy: 0.8337
Epoch 179/200
185/185 [=====] - 7s 38ms/step - loss: 0.1709 - accuracy: 0.9471 -
val_loss: 0.7598 - val_accuracy: 0.8337
Epoch 180/200
185/185 [=====] - 7s 38ms/step - loss: 0.1682 - accuracy: 0.9478 -
val_loss: 0.7596 - val_accuracy: 0.8433
Epoch 181/200
185/185 [=====] - 7s 38ms/step - loss: 0.1488 - accuracy: 0.9534 -
val_loss: 0.8411 - val_accuracy: 0.8260
Epoch 182/200
185/185 [=====] - 7s 38ms/step - loss: 0.1531 - accuracy: 0.9542 -
val_loss: 0.7745 - val_accuracy: 0.8423
Epoch 183/200
185/185 [=====] - 7s 38ms/step - loss: 0.1608 - accuracy: 0.9493 -
val_loss: 0.7912 - val_accuracy: 0.8173
Epoch 184/200
185/185 [=====] - 8s 41ms/step - loss: 0.1677 - accuracy: 0.9507 -
val_loss: 0.8367 - val_accuracy: 0.8154
Epoch 185/200
185/185 [=====] - 8s 41ms/step - loss: 0.1516 - accuracy: 0.9517 -
val_loss: 0.7287 - val_accuracy: 0.8346
Epoch 186/200
185/185 [=====] - 8s 41ms/step - loss: 0.1418 - accuracy: 0.9568 -
val_loss: 0.8128 - val_accuracy: 0.8096
Epoch 187/200
185/185 [=====] - 7s 38ms/step - loss: 0.1566 - accuracy: 0.9490 -
val_loss: 0.8848 - val_accuracy: 0.8106
Epoch 188/200
185/185 [=====] - 7s 38ms/step - loss: 0.2359 - accuracy: 0.9293 -
val_loss: 0.7993 - val_accuracy: 0.8125
Epoch 189/200
185/185 [=====] - 7s 38ms/step - loss: 0.1763 - accuracy: 0.9437 -
val_loss: 0.7120 - val_accuracy: 0.8365
Epoch 190/200
185/185 [=====] - 7s 38ms/step - loss: 0.1622 - accuracy: 0.9496 -
val_loss: 0.7484 - val_accuracy: 0.8327
Epoch 191/200
185/185 [=====] - 7s 38ms/step - loss: 0.1437 - accuracy: 0.9544 -
val_loss: 0.7308 - val_accuracy: 0.8385
Epoch 192/200
185/185 [=====] - 7s 38ms/step - loss: 0.1481 - accuracy: 0.9529 -
val_loss: 0.8170 - val_accuracy: 0.8231
Epoch 193/200
185/185 [=====] - 7s 38ms/step - loss: 0.1295 - accuracy: 0.9585 -
val_loss: 0.9294 - val_accuracy: 0.8067
Epoch 194/200
185/185 [=====] - 7s 38ms/step - loss: 0.1294 - accuracy: 0.9619 -
val_loss: 0.7718 - val_accuracy: 0.8327
Epoch 195/200
```

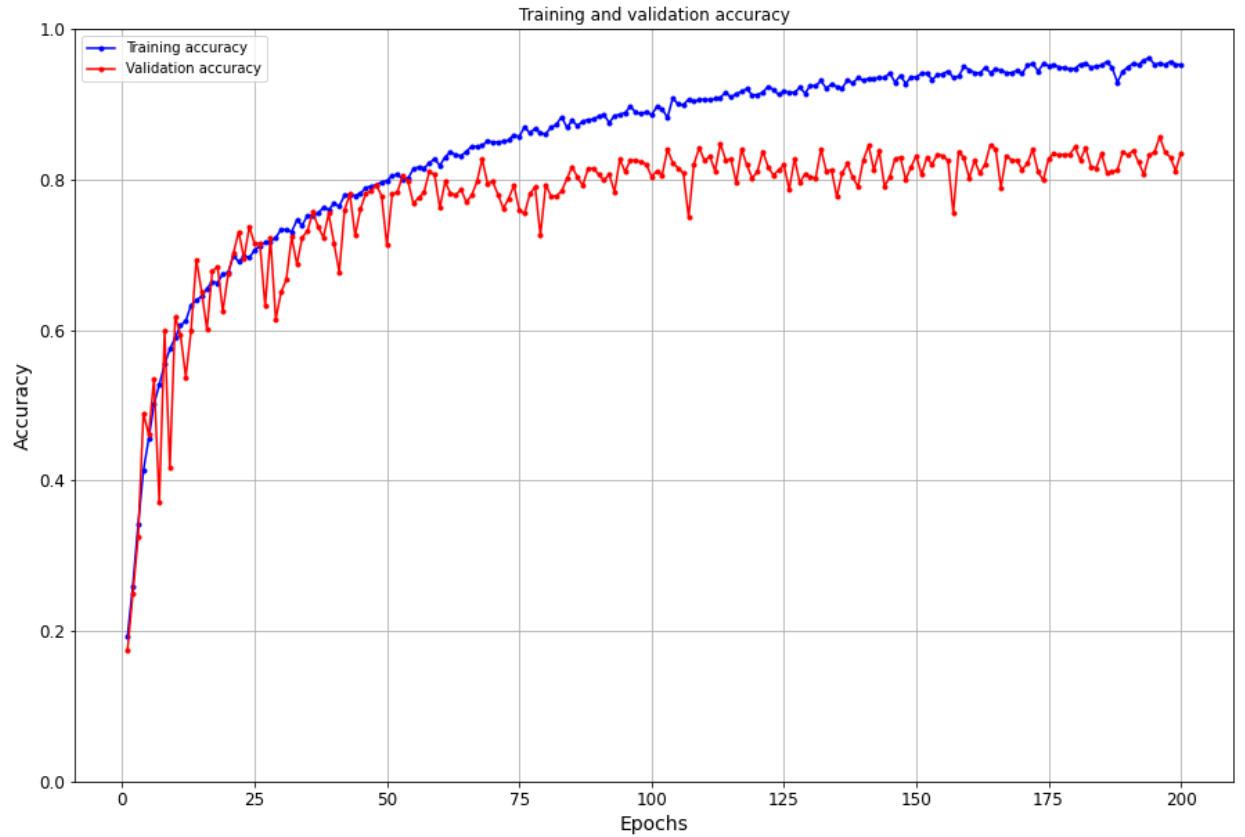
```
185/185 [=====] - 7s 38ms/step - loss: 0.1501 - accuracy: 0.9530 -  
val_loss: 0.8341 - val_accuracy: 0.8365  
Epoch 196/200  
185/185 [=====] - 7s 38ms/step - loss: 0.1411 - accuracy: 0.9549 -  
val_loss: 0.7799 - val_accuracy: 0.8567  
Epoch 197/200  
185/185 [=====] - 7s 38ms/step - loss: 0.1561 - accuracy: 0.9527 -  
val_loss: 0.7753 - val_accuracy: 0.8365  
Epoch 198/200  
185/185 [=====] - 7s 38ms/step - loss: 0.1474 - accuracy: 0.9569 -  
val_loss: 0.8242 - val_accuracy: 0.8298  
Epoch 199/200  
185/185 [=====] - 7s 38ms/step - loss: 0.1464 - accuracy: 0.9534 -  
val_loss: 0.8546 - val_accuracy: 0.8115  
Epoch 200/200  
185/185 [=====] - 7s 38ms/step - loss: 0.1477 - accuracy: 0.9527 -  
val_loss: 0.7284 - val_accuracy: 0.8356
```

```
In [ ]: plot_loss_and_accuracy(history)
```

```
accuracy      0.961851  
dtype: float64 val_accuracy      0.856731  
dtype: float64
```



```
<Figure size 432x288 with 0 Axes>
```



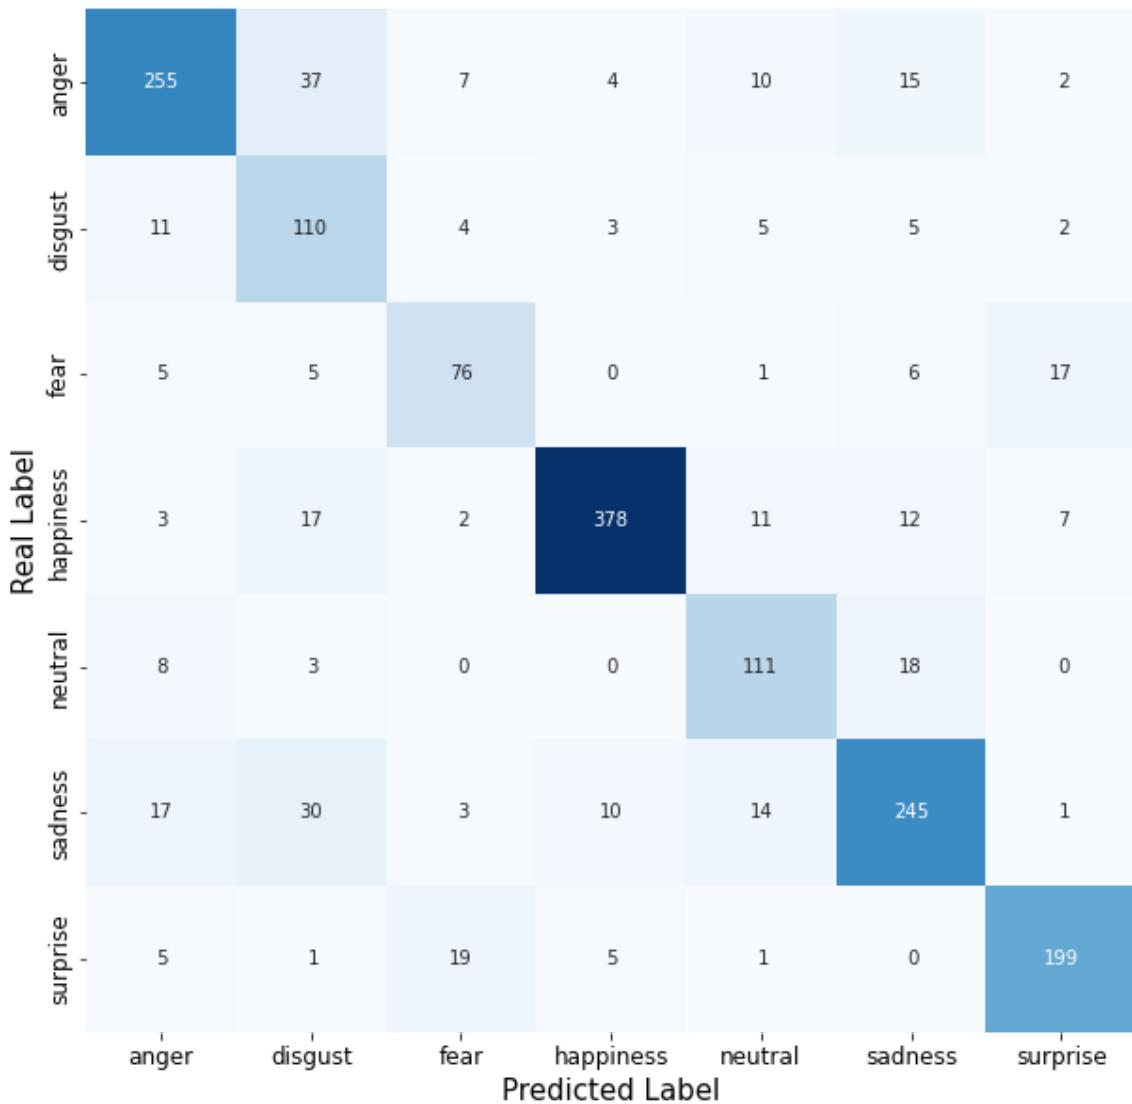
```
In [ ]: modelWithAugmentation.evaluate(test_ds)
```

```
54/54 [=====] - 1s 16ms/step - loss: 0.9021 - accuracy: 0.8082
```

```
Out[ ]: [0.9020760655403137, 0.8082352876663208]
```

```
In [ ]: y_pred = modelWithAugmentation.predict(test_ds)
cnf_matrix = plot_confusion_matrix(y_pred, y_test, class_names, return_m=True)
```

Confusion Matrix



```
In [ ]: print_confusion_matrix_metrics(cnf_matrix)
```

False Positives:

```
anger -> 49
disgust -> 93
fear -> 35
happiness -> 22
neutral -> 42
sadness -> 56
surprise -> 29
```

False Negatives:

```
anger -> 75
disgust -> 30
fear -> 34
happiness -> 52
neutral -> 29
sadness -> 75
surprise -> 31
```

True Positives:

```
anger -> 255
disgust -> 110
fear -> 76
happiness -> 378
neutral -> 111
sadness -> 245
```

surprise -> 199

True Negatives:

anger -> 1321
disgust -> 1467
fear -> 1555
happiness -> 1248
neutral -> 1518
sadness -> 1324
surprise -> 1441

True Positive Rate:

anger -> 0.773%
disgust -> 0.786%
fear -> 0.691%
happiness -> 0.879%
neutral -> 0.793%
sadness -> 0.766%
surprise -> 0.865%

True Negative Rate:

anger -> 0.964%
disgust -> 0.940%
fear -> 0.978%
happiness -> 0.983%
neutral -> 0.973%
sadness -> 0.959%
surprise -> 0.980%

Positive Predictive Value:

anger -> 1
disgust -> 1
fear -> 1
happiness -> 1
neutral -> 1
sadness -> 1
surprise -> 1

Negative Predictive Value:

anger -> 1
disgust -> 1
fear -> 1
happiness -> 1
neutral -> 1
sadness -> 1
surprise -> 1

False Positive Rate:

anger -> 0.036%
disgust -> 0.060%
fear -> 0.022%
happiness -> 0.017%
neutral -> 0.027%
sadness -> 0.041%
surprise -> 0.020%

False Negative Rate:

anger -> 0.227%
disgust -> 0.214%
fear -> 0.309%
happiness -> 0.121%
neutral -> 0.207%
sadness -> 0.234%
surprise -> 0.135%

False Discovery Rate:

```
anger -> 0.161%
disgust -> 0.458%
fear -> 0.315%
happiness -> 0.055%
neutral -> 0.275%
sadness -> 0.186%
surprise -> 0.127%
```

Accuracy:

```
anger -> 0.927%
disgust -> 0.928%
fear -> 0.959%
happiness -> 0.956%
neutral -> 0.958%
sadness -> 0.923%
surprise -> 0.965%
```

```
In [ ]: print_general_metrics(y_pred, y_test)
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.77	0.80	330
1	0.54	0.79	0.64	140
2	0.68	0.69	0.69	110
3	0.94	0.88	0.91	430
4	0.73	0.79	0.76	140
5	0.81	0.77	0.79	320
6	0.87	0.87	0.87	230
accuracy			0.81	1700
macro avg	0.77	0.79	0.78	1700
weighted avg	0.82	0.81	0.81	1700

Accuracy: 0.808

Balanced accuracy: 0.793

F1_Score: 0.812

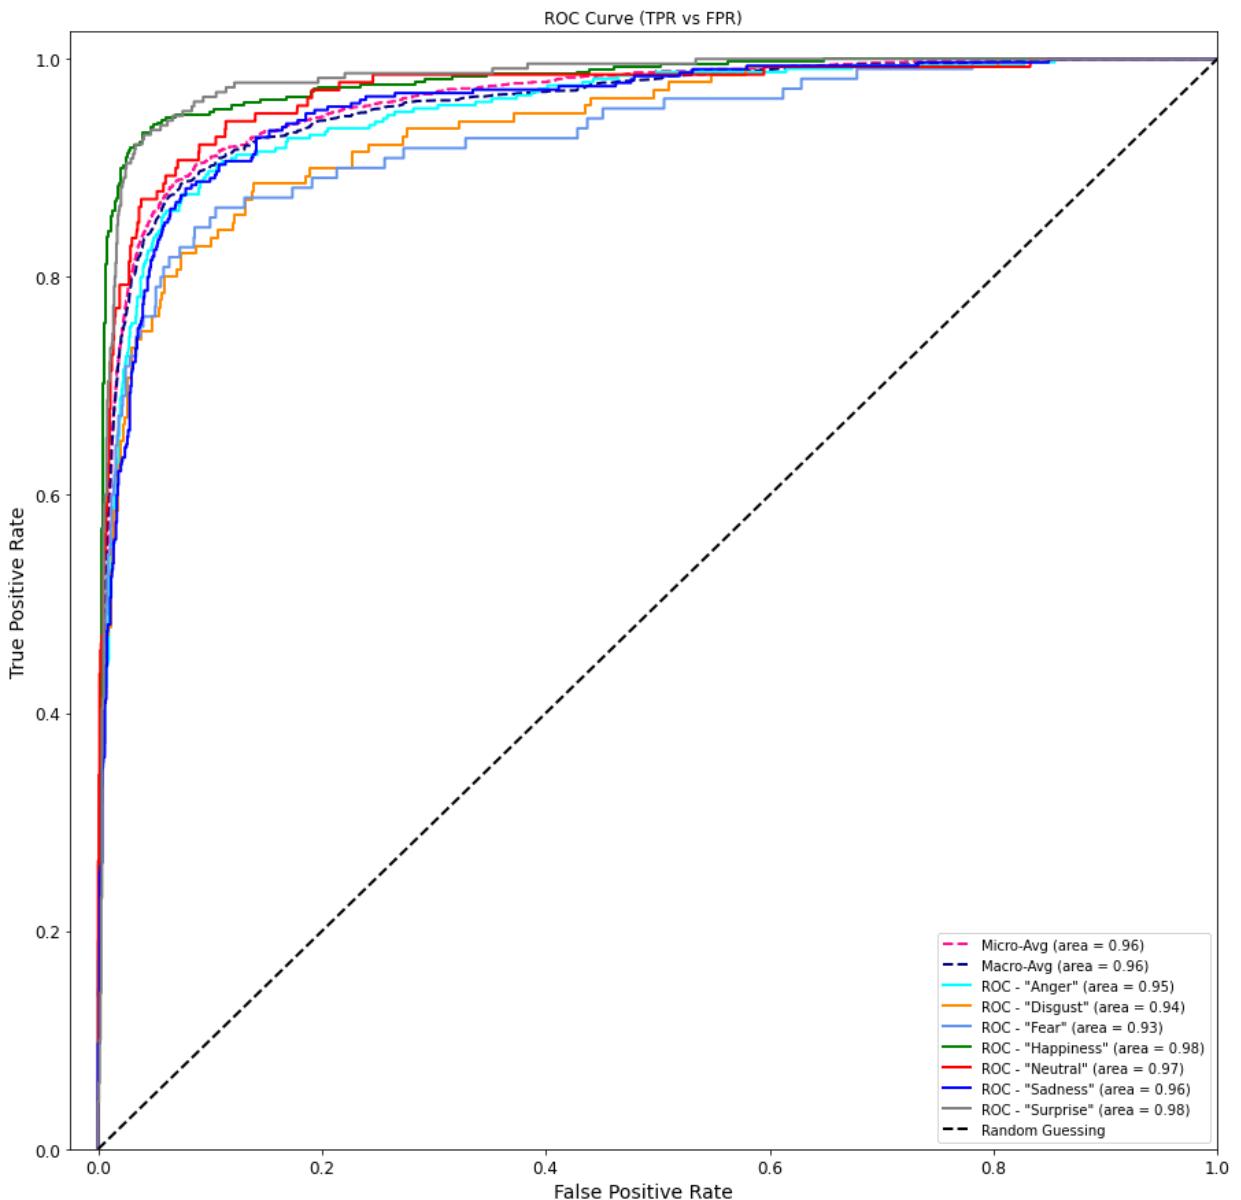
Precision: 0.822

Recall: 0.808

ROC AUC (OVR): 0.959

ROC AUC (OVO): 0.958

```
In [ ]: plot_multiclass_roc_curve(y_test, y_pred, class_names)
```



One-vs-One ROC AUC scores:

0.958685 (macro),
0.963208 (weighted by prevalence)

One-vs-Rest ROC AUC scores:

0.958685 (macro),
0.963208 (weighted by prevalence)

In []:

```
# Start TensorBoard
%load_ext tensorboard
# --port=6006
# --host=127.0.0.1
%tensorboard --logdir=./tensorboard_logs_CNN_V1_Rafael --host=127.0.0.1
```

The tensorboard extension is already loaded. To reload it, use:

%reload_ext tensorboard

Reusing TensorBoard on port 6006 (pid 24208), started 1:03:38 ago. (Use '!kill 24208' to kill it.)

Apply trained model to new images

In []:

```
from keras.models import load_model
model = load_model(best_model_file_path)
model.evaluate(test_ds)
```

54/54 [=====] - 1s 10ms/step - loss: 0.8130 - accuracy: 0.8382

Out[]: [0.8130242824554443, 0.8382353186607361]

In []:

```
from keras.preprocessing import image

def predict_new_image(image_path, model, class_names):
    img = image.load_img(image_path, target_size=(80,80))
    img_as_array = image.img_to_array(img)
    plt.imshow( img_as_array / 255)
```

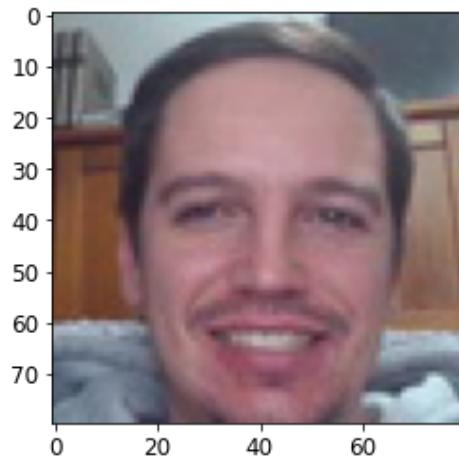
```
x=img_as_array
x= x/ 255.
x=np.expand_dims(x, axis=0)
images = np.vstack([x])

classes = model.predict(images, batch_size=16)

print(classes.argmax(axis=-1))
print(class_names)
print(np.array(class_names)[classes.argmax(axis=-1)])
```

In []: predict_new_image("im1.jpg", model, class_names)

[3]
['anger', 'disgust', 'fear', 'happiness', 'neutral', 'sadness', 'surprise']
['happiness']



In []: predict_new_image("im2.jpg", model, class_names)

[5]
['anger', 'disgust', 'fear', 'happiness', 'neutral', 'sadness', 'surprise']
['sadness']

